



CLOUD COMPUTING CONCEPTS

with Indranil Gupta (Indy)

P2P SYSTEMS

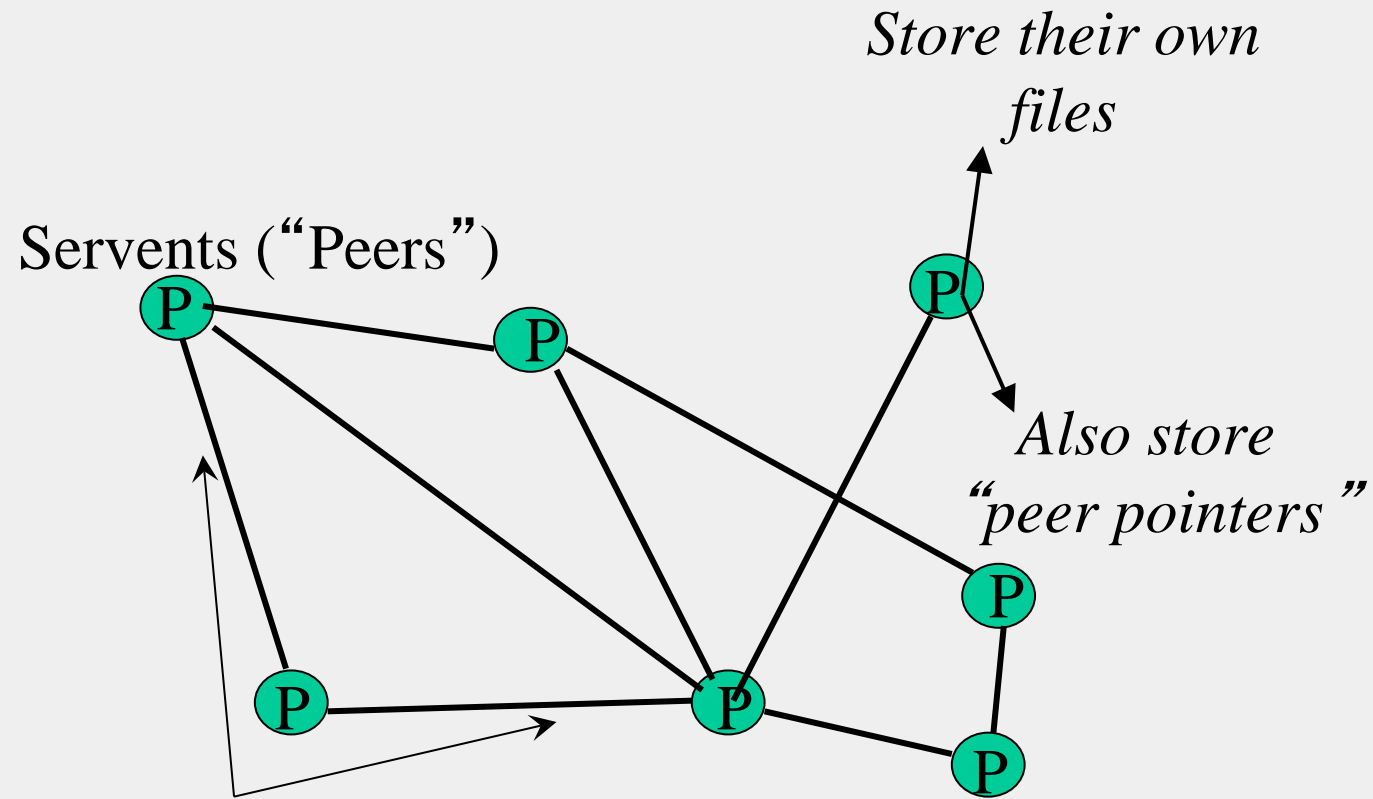
Lecture C

GNUTELLA

GNUTELLA

- Eliminate the servers
- Client machines search and retrieve amongst themselves
- Clients act as servers too, called **servents**
- [3/00] release by AOL, immediately withdrawn, but 88K users by 3/03
- Original design underwent several modifications

GNUTELLA

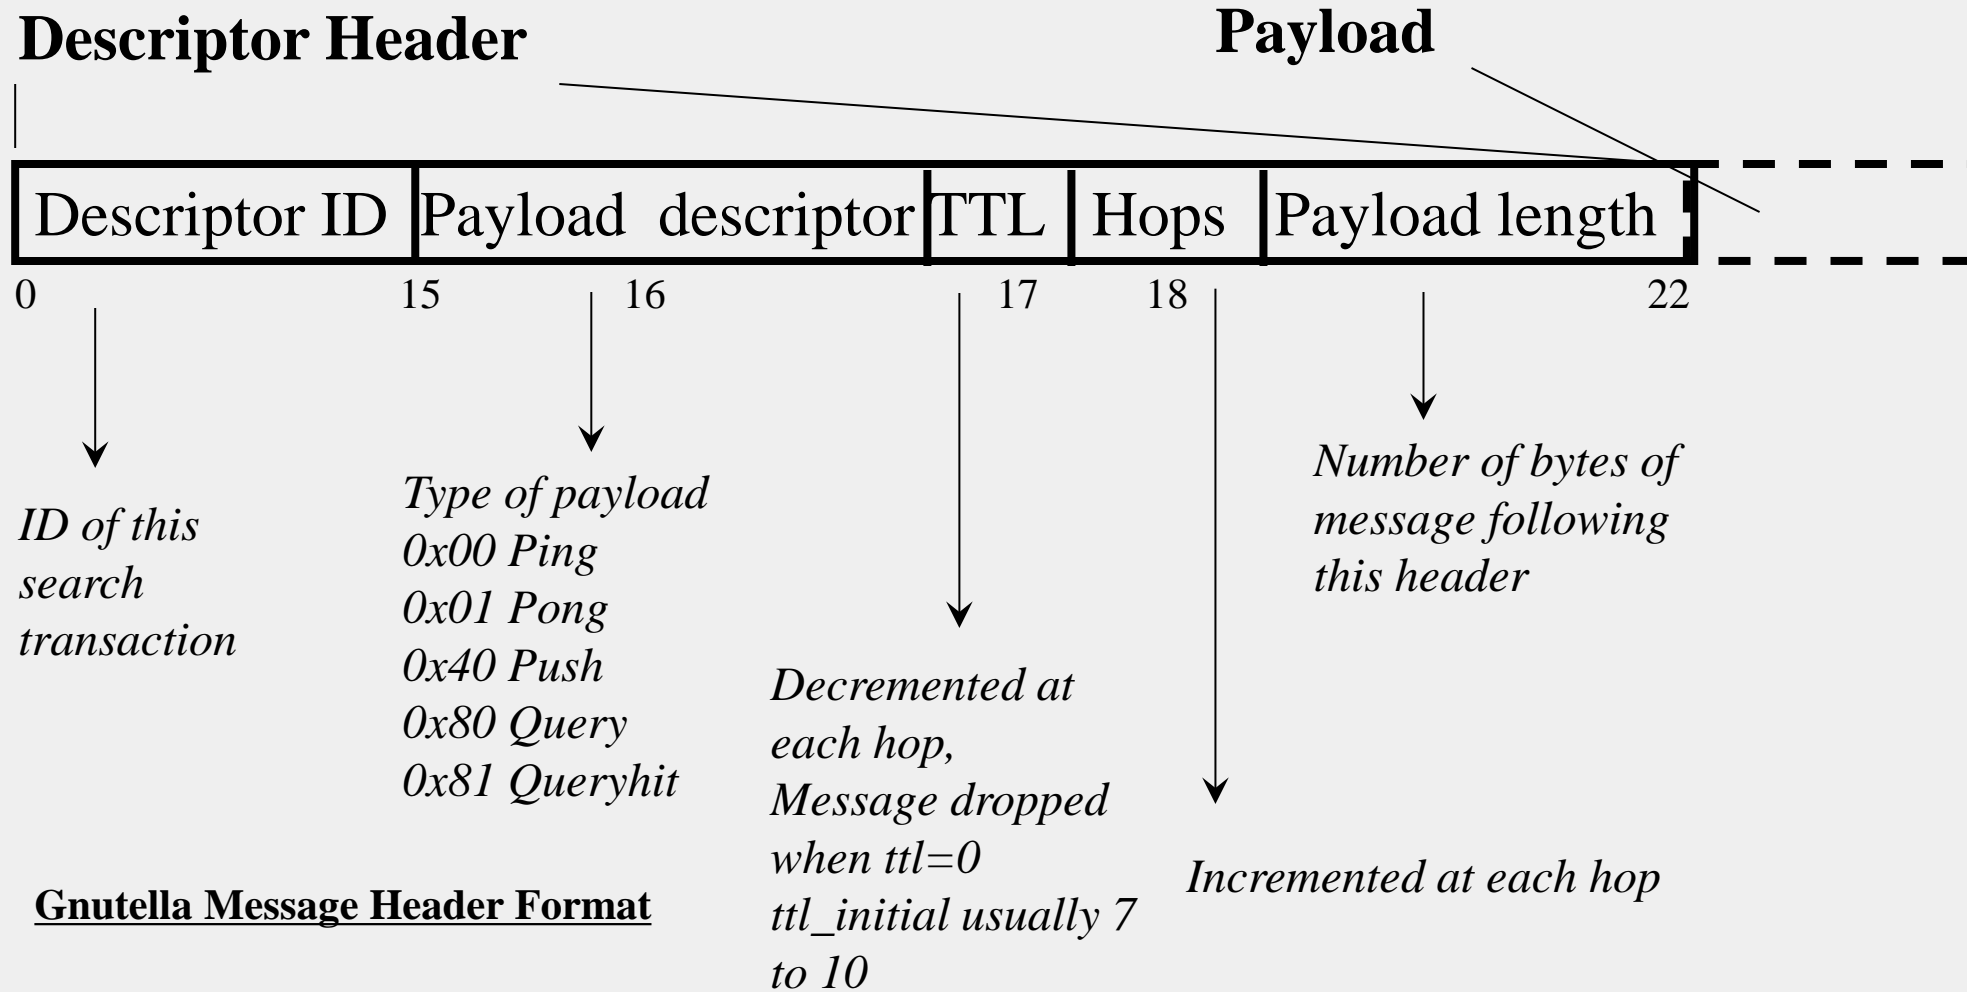


Connected in an **overlay graph**
(== each link is an implicit Internet path)

How do I search for my BEATLES file?


- Gnutella *routes* different messages within the overlay graph
- Gnutella protocol has 5 main message types
 - **Query** (search)
 - **QueryHit** (response to query)
 - **Ping** (to probe network for other peers)
 - **Pong** (reply to ping, contains address of another peer)
 - Push (used to initiate file transfer)
- We'll go into the message structure and protocol now
 - All fields except IP address are in little-endian format
 - 0x12345678 stored as 0x78 in lowest address byte, then 0x56 in next higher address, and so on.

How do I SEARCH FOR MY BEATLES FILE?



How do I SEARCH FOR MY BEATLES FILE?

Query (0x80)



Payload Format in Gnutella **Query** Message

TTL=2

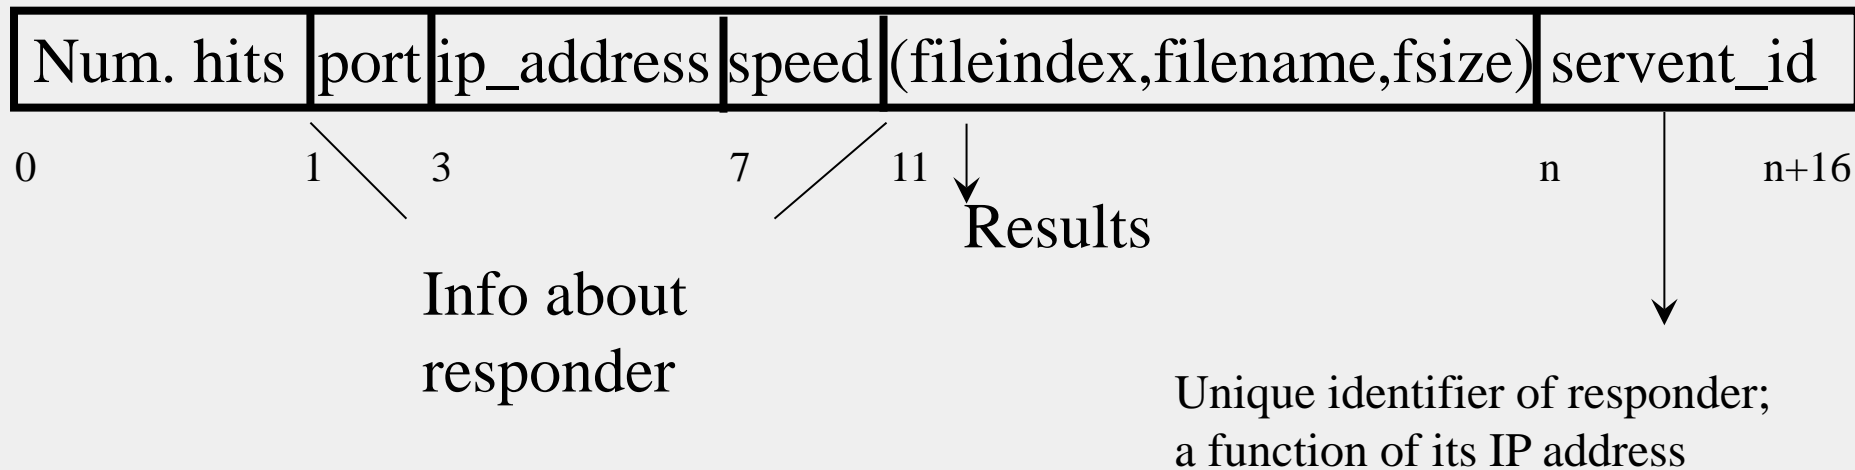
Who has PennyLane.mp3?

TTL=2

Who has PennyLane.mp3?

GNUTELLA SEARCH


QueryHit (0x81) : successful result to a query



Payload Format in Gnutella **QueryHit** Message

GNUTELLA SEARCH

Successful results **QueryHit**'s routed on reverse path



AVOIDING EXCESSIVE TRAFFIC

- To avoid duplicate transmissions, each peer maintains a list of recently received messages
- Query forwarded to all neighbors except peer from which received
- Each Query (identified by DescriptorID) forwarded only once
- QueryHit routed back only to peer from which Query received with same DescriptorID
- Duplicates with same DescriptorID and Payload descriptor (msg type) are dropped
- QueryHit with DescriptorID for which Query not seen is dropped

AFTER RECEIVING QUERY HIT MESSAGES

- Requestor chooses “best” QueryHit responder
 - Initiates HTTP request directly to responder’s ip+port

```
GET /get/<File Index>/<File Name>/HTTP/1.0\r\n
```

```
Connection: Keep-Alive\r\n
```

```
Range: bytes=0-\r\n
```

```
User-Agent: Gnutella\r\n
```

```
\r\n
```

- Responder then replies with file packets after this message:

```
HTTP 200 OK\r\n
```

```
Server: Gnutella\r\n
```

```
Content-type:application/binary\r\n
```

```
Content-length: 1024 \r\n
```


```
\r\n
```

AFTER RECEIVING QUERY HIT MESSAGES (2)

- HTTP is the file transfer protocol. Why?
 - Because it's standard, well-debugged, and widely used.
- Why the “range” field in the GET request?
 - To support partial file transfers.
- What if responder is behind firewall that disallows incoming connections?

DEALING WITH FIREWALLS

Requestor sends **Push** to responder asking for file transfer



DEALING WITH FIREWALLS

Push (0x40)



same as in
received QueryHit

Address at which
requestor can accept
incoming connections

DEALING WITH FIREWALLS

- Responder establishes a TCP connection at ip_address, port specified. Sends

GIV <File Index>:<Servent Identifier>/<File Name>\n\n

- Requestor then sends GET to responder (as before) and file is transferred as explained earlier
- What if requestor is behind firewall too?
 - Gnutella gives up
 - Can you think of an alternative solution?

PING-PONG

Ping (0x00)

no payload

Pong (0x01)

Port	ip_address	Num. files shared	Num. KB shared
------	------------	-------------------	----------------

- Peers initiate Ping's periodically
- Ping's flooded out like Query's, Pong's routed along reverse path like QueryHit's
- Pong replies used to update set of neighboring peers
 - To keep neighbor lists fresh in spite of peers joining, leaving and failing

GNUTELLA SUMMARY

- No servers
- Peers/servents maintain “neighbors,” this forms an overlay graph
- Peers store their own files
- Queries flooded out, ttl restricted
- QueryHit (replies) reverse path routed
- Supports file transfer through firewalls
- Periodic ping-pong to continuously refresh neighbor lists
 - List size specified by user at peer: heterogeneity means some peers may have more neighbors
 - Gnutella found to follow **power law** distribution:
$$P(\text{\#links} = L) \sim L^{-k} \quad (k \text{ is a constant})$$

PROBLEMS

- Ping/Pong constituted 50% traffic
 - Solution: Multiplex, *cache* and reduce frequency of pings/pongs
- Repeated searches with same keywords
 - Solution: *Cache* Query, QueryHit messages
- Modem-connected hosts do not have enough bandwidth for passing Gnutella traffic
 - Solution: use a central server to act as proxy for such peers
 - Another solution:
 - ➔ FastTrack System (soon)

PROBLEMS (CONTD.)

- Large number of *freeloaders*
 - 70% of users in 2000 were freeloaders
 - Only download files, never upload own files
- Flooding causes excessive traffic
 - Is there some way of maintaining meta-information about peers that leads to more intelligent routing?
 - ➔ Structured peer-to-peer systems
e.g., Chord System (coming up soon)