



# CLOUD COMPUTING CONCEPTS

with Indranil Gupta (Indy)

## P2P SYSTEMS

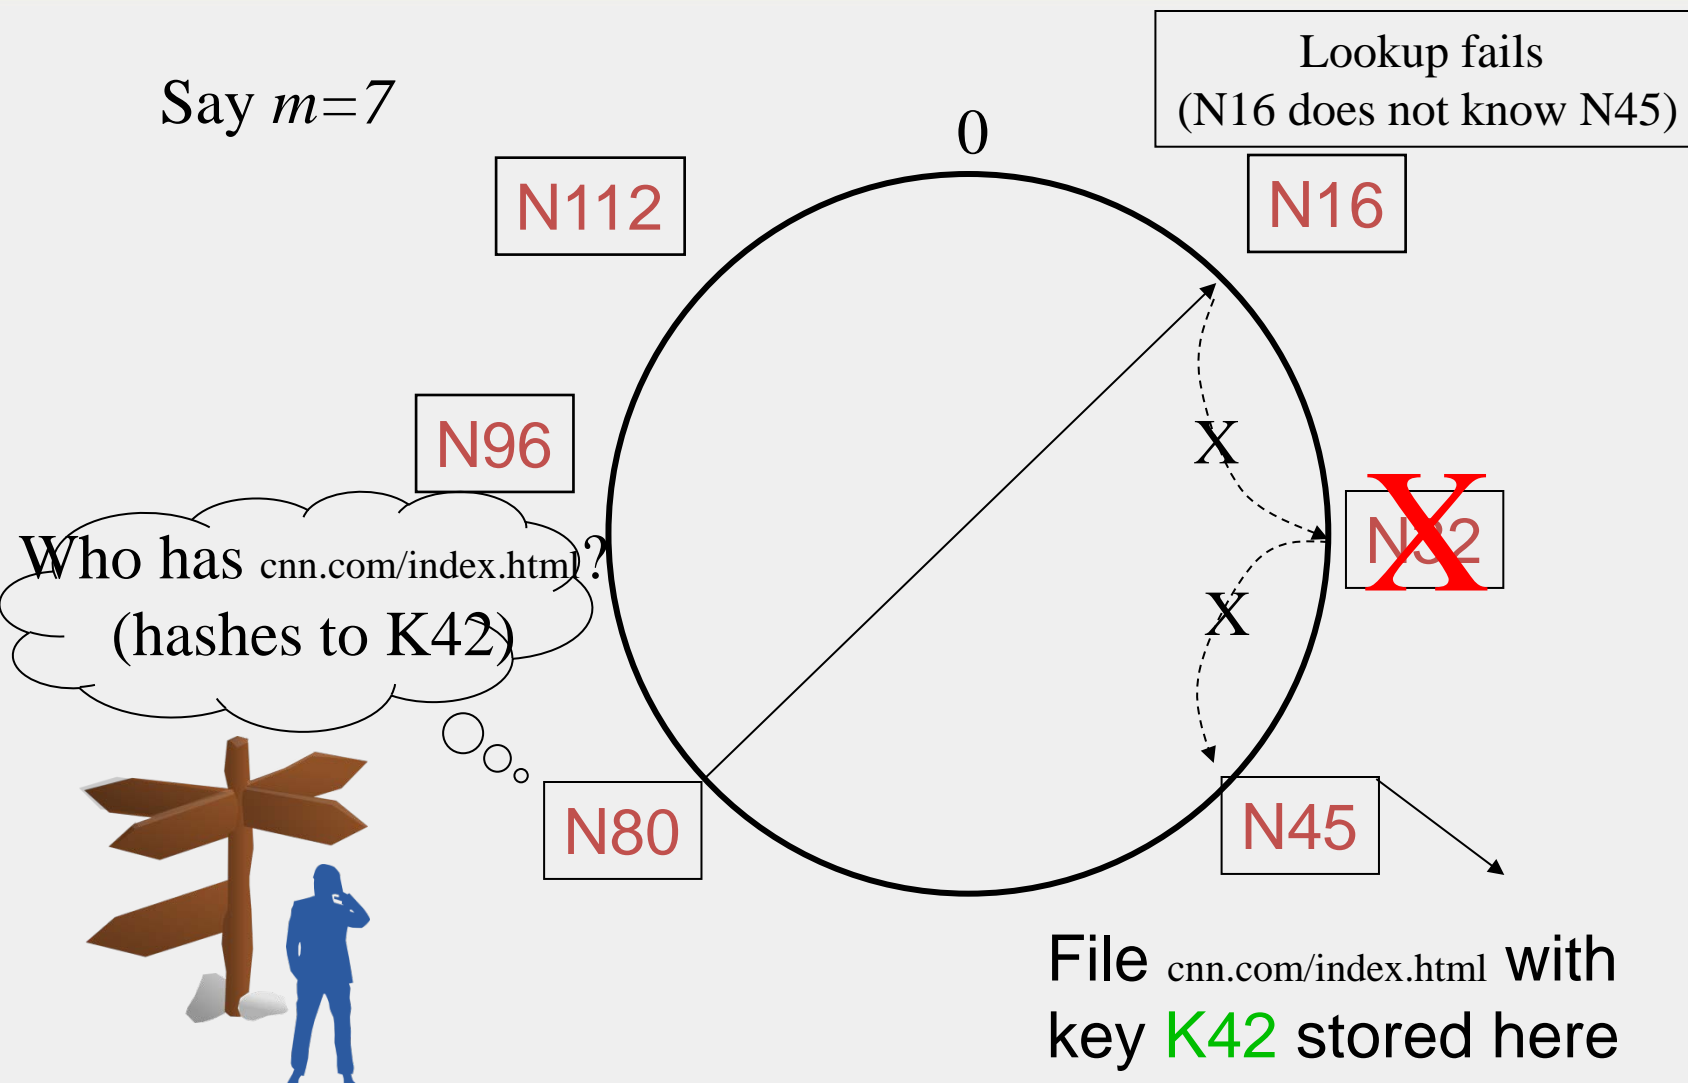
Lecture F

---

### FAILURES IN CHORD

# SEARCH UNDER PEER FAILURES


Say  $m=7$



# SEARCH UNDER PEER FAILURES

Say  $m=7$

One solution: maintain  $r$  multiple *successor* entries  
0 In case of failure, use successor entries



# SEARCH UNDER PEER FAILURES

- Choosing  $r=2\log(N)$  suffices to maintain *lookup correctness* w.h.p. (i.e., ring connected)
  - Say 50% of nodes fail
  - $\Pr(\text{at given node, at least one successor alive})=$


$$1 - \left(\frac{1}{2}\right)^{2\log N} = 1 - \frac{1}{N^2}$$

- $\Pr(\text{above is true at all alive nodes})=$

$$\left(1 - \frac{1}{N^2}\right)^{N/2} = e^{-\frac{1}{2N}} \approx 1$$

# SEARCH UNDER PEER FAILURES (2)

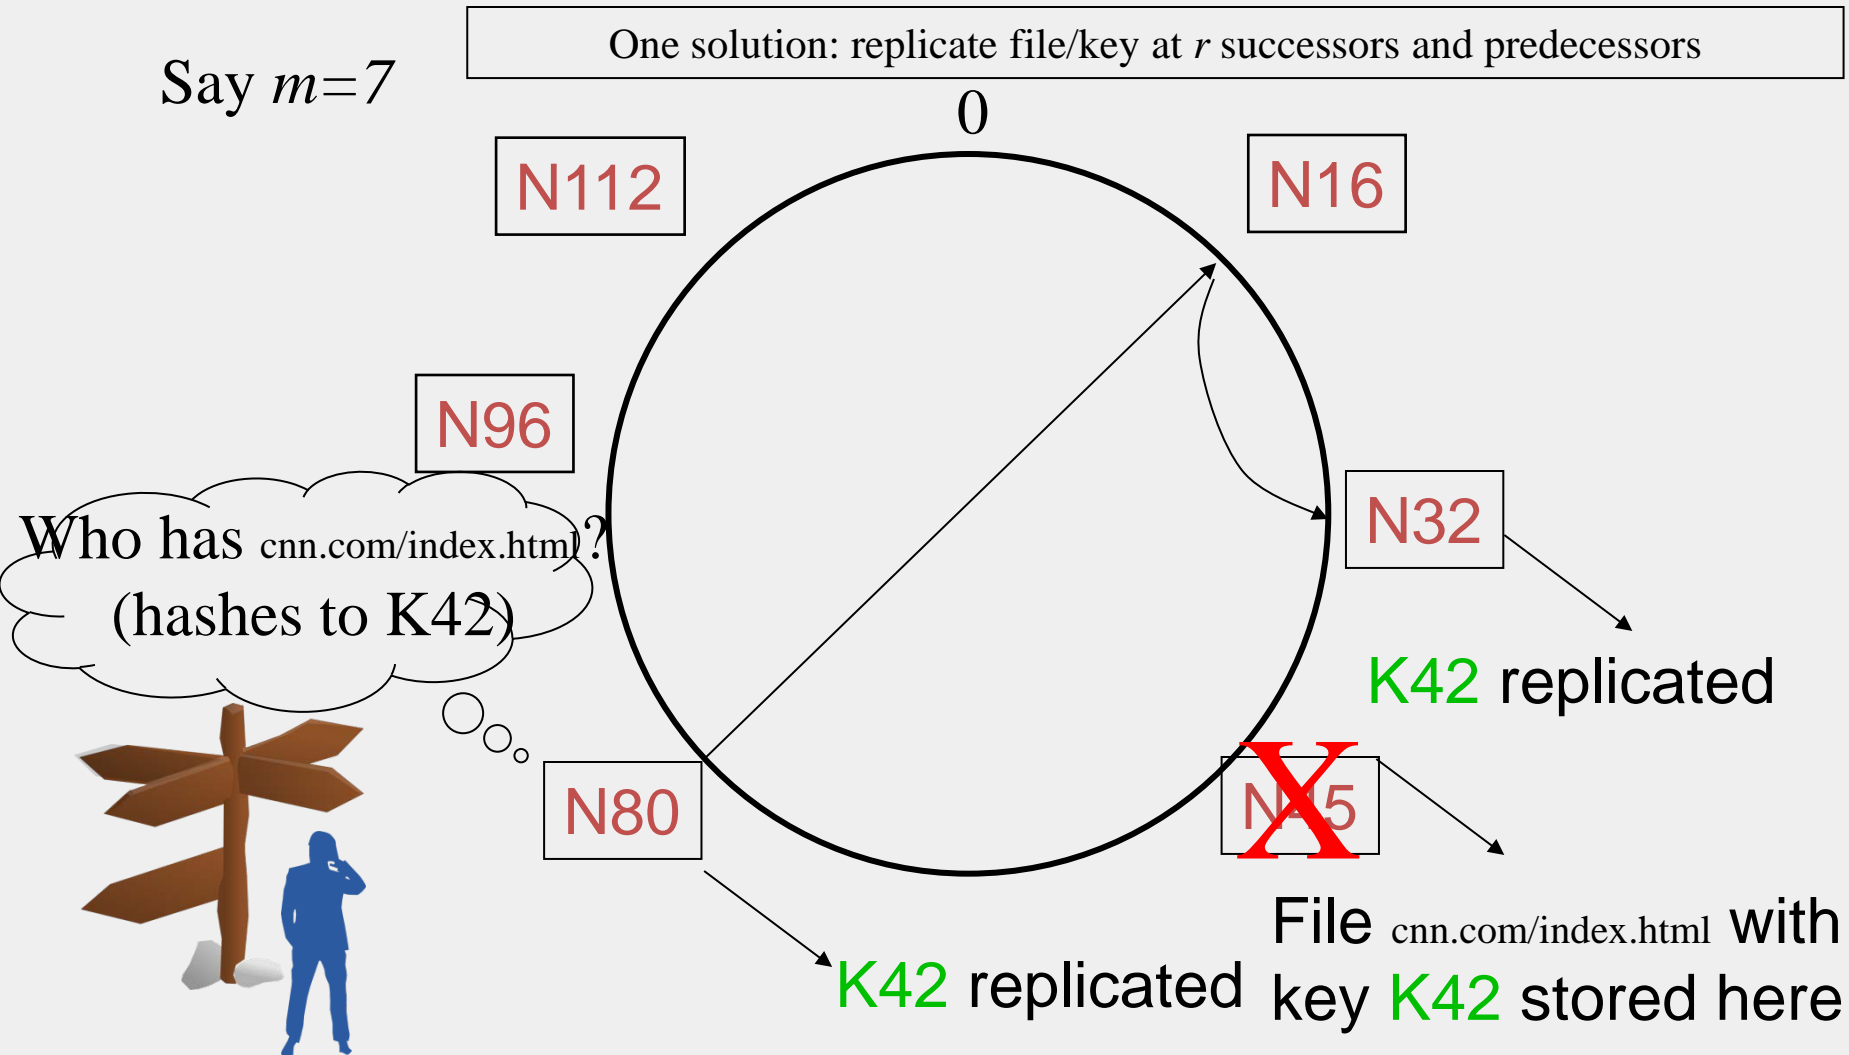
Say  $m=7$



# SEARCH UNDER PEER FAILURES (2)

Say  $m=7$

One solution: replicate file/key at  $r$  successors and predecessors



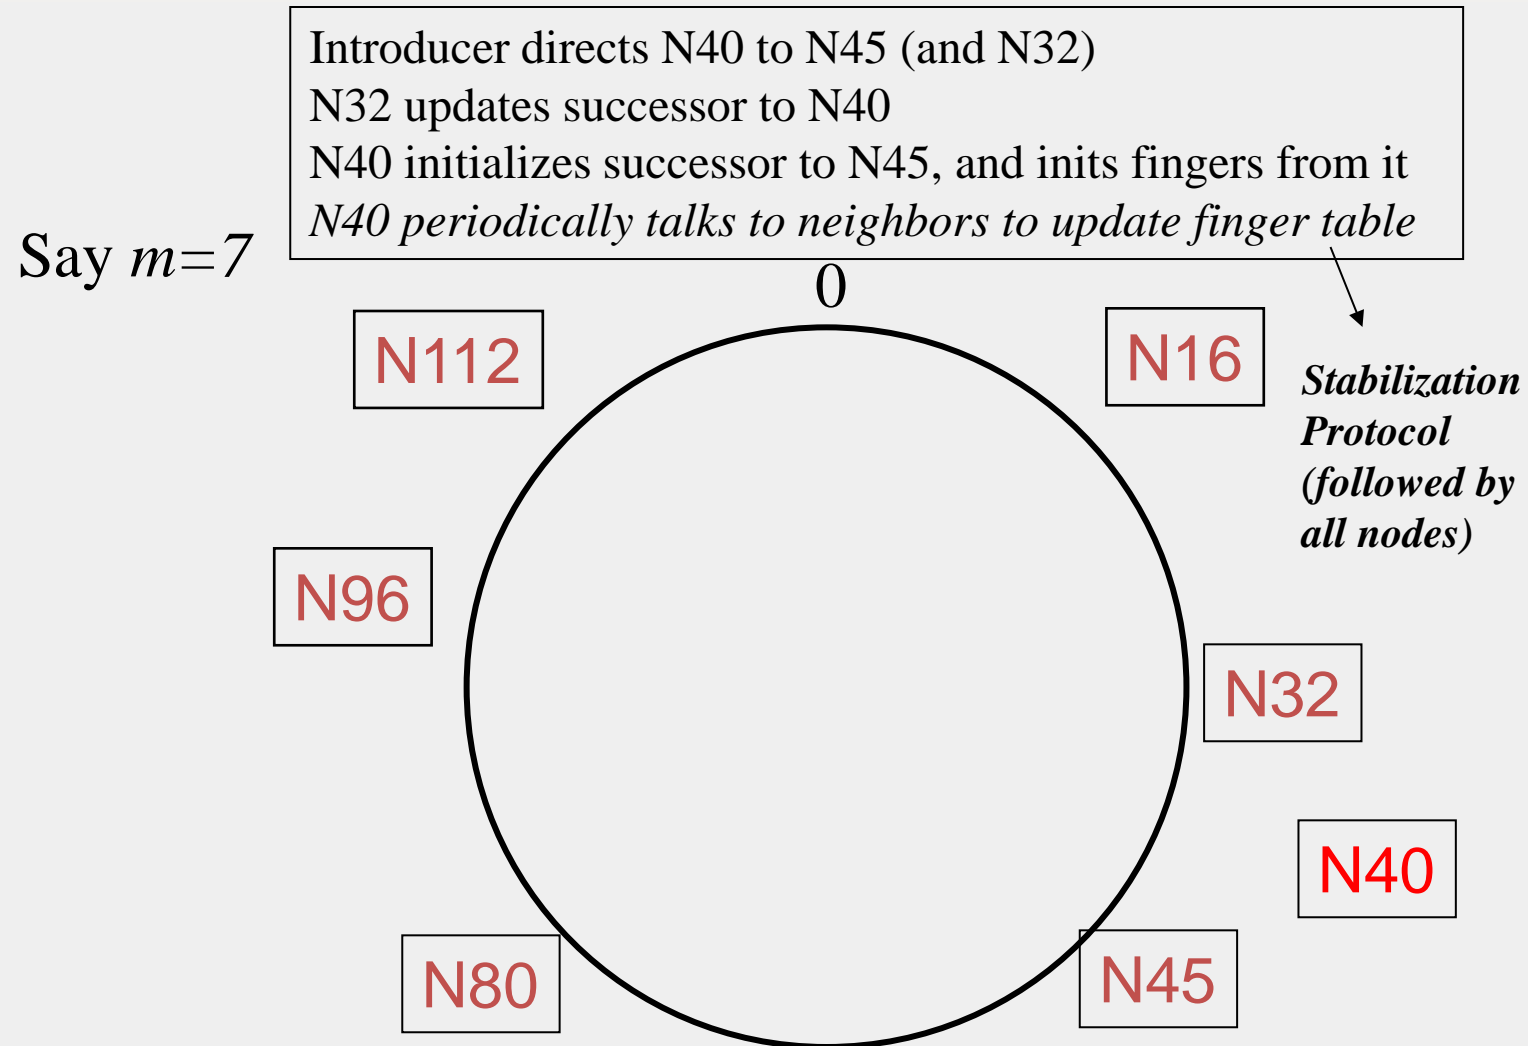
# NEED TO DEAL WITH DYNAMIC CHANGES

- ✓ Peers fail
- New peers join
- Peers leave
  - P2P systems have a high rate of *churn* (node join, leave and failure)
    - 25% per hour in Overnet (eDonkey)
    - 100% per hour in Gnutella
    - Lower in managed clusters
    - Common feature in all distributed systems, including wide-area (e.g., PlanetLab), clusters (e.g., Emulab), clouds (e.g., AWS), etc.

So, all the time, need to:

→ update *successors* and *fingers*, and copy keys

# NEW PEERS JOINING






# NEW PEERS JOINING (2)

N40 may need to copy some files/keys from N45  
(files with fileid between 32 and 40)

Say  $m=7$



## NEW PEERS JOINING (3)

- A new peer affects  $O(\log(N))$  other finger entries in the system, on average [Why?]
- Number of messages per peer join =  $O(\log(N) * \log(N))$
- Similar set of operations for dealing with peers leaving
  - For dealing with failures, also need *failure detectors* (we'll see these later in the course!)

# STABILIZATION PROTOCOL

- Concurrent peer joins, leaves, failures might cause loopiness of pointers and failure of lookups
  - Chord peers periodically run a *stabilization* algorithm that checks and updates pointers and keys
  - Ensures *non-loopiness* of fingers, eventual success of lookups and  $O(\log(N))$  lookups w.h.p.
  - Each stabilization round at a peer involves a constant number of messages
  - Strong stability takes  $O(N^2)$  stabilization rounds
  - For more see [TechReport on Chord webpage]

# CHURN

- When nodes are constantly joining, leaving, failing
  - Significant effect to consider: traces from the Overnet system show *hourly* peer turnover rates (*churn*) could be 25–100% of total number of nodes in system
  - Leads to excessive (unnecessary) key copying (remember that keys are replicated)
  - Stabilization algorithm may need to consume more bandwidth to keep up
  - Main issue is that files are replicated, while it might be sufficient to replicate only meta information about files
  - Alternatives
    - Introduce a level of indirection (any p2p system)
    - Replicate metadata more, e.g., Kelips (later in this lecture series)

# VIRTUAL NODES

- Hash can get non-uniform → Bad load balancing
  - Treat each node as multiple virtual nodes behaving independently
  - Each joins the system
  - Reduces variance of load imbalance

# WRAP-UP NOTES

- Virtual Ring and Consistent Hashing used in Cassandra, Riak, Voldemort, DynamoDB, and other key-value stores
- Current status of Chord project:
  - File systems (CFS, Ivy) built on top of Chord
  - DNS lookup service built on top of Chord
  - Internet Indirection Infrastructure (I3) project at UC Berkeley
  - Spawned research on many interesting issues about p2p systems

<http://www.pdos.lcs.mit.edu/chord/>