



# CLOUD COMPUTING CONCEPTS

---

with **Indranil Gupta (Indy)**

## KEY-VALUE STORES NoSQL

Lecture A

---

### WHY KEY-VALUE/NoSQL?

# THE KEY-VALUE ABSTRACTION

- (Business) Key  $\rightarrow$  Value
- (twitter.com) Tweet id  $\rightarrow$  information about tweet
- (amazon.com) Item number  $\rightarrow$  information about it
- (kayak.com) Flight number  $\rightarrow$  information about flight, e.g., availability
- (yourbank.com) Account number  $\rightarrow$  information about it



# THE KEY-VALUE ABSTRACTION (2)

- It's a dictionary datastructure.
  - Insert, lookup, and delete by key
  - E.g., hash table, binary tree
- But distributed
- Sound familiar? Remember distributed hash tables (DHT) in P2P systems?
- It's not surprising that key-value stores reuse many techniques from DHTs.



# ISN'T THAT JUST A DATABASE?

- Yes, sort of
- Relational Database Management Systems (RDBMSs) have been around for ages
- MySQL is the most popular among them
- Data stored in tables
- Schema-based, i.e., structured tables
- Each row (data item) in a table has a primary key that is unique within that table
- Queried using SQL (Structured Query Language)
- Supports joins



# RELATIONAL DATABASE EXAMPLE

**users table**

user_id	name	zipcode	blog_url	blog_id
101	Alice	12345	alice.net	1
422	Charlie	45783	charlie.com	3
555	Bob	99910	bob.blogspot.com	2

↑  
Primary keys

↑  
Foreign keys

**blog table**

id	url	last_updated	num_posts
1	alice.net	5/2/14	332
2	bob.blogspot.com	4/2/13	10003
3	charlie.com	6/15/14	7

## Example SQL queries

1. `SELECT zipcode  
FROM users  
WHERE name = "Bob"`
2. `SELECT url  
FROM blog  
WHERE id = 3`
3. `SELECT users.zipcode, blog.num_posts  
FROM users JOIN blog  
ON users.blog_url = blog.url`



# MISMATCH WITH TODAY'S WORKLOADS

- Data: Large and unstructured
- Lots of random reads and writes
- Sometimes write-heavy
- Foreign keys rarely needed
- Joins infrequent



# NEEDS OF TODAY'S WORKLOADS

- Speed
- Avoid Single Point of Failure (SPOF)
- Low TCO (Total cost of operation)
- Fewer system administrators
- Incremental scalability
- Scale out, not up
  - What?



# SCALE OUT, NOT SCALE UP

- Scale up = grow your cluster capacity by replacing with more powerful machines
  - Traditional approach
  - Not cost-effective, as you're buying above the sweet spot on the price curve
  - And you need to replace machines often
- Scale out = incrementally grow your cluster capacity by adding more COTS machines (Components Off the Shelf)
  - Cheaper
  - Over a long duration, phase in a few newer (faster) machines as you phase out a few older machines
  - Used by most companies who run datacenters and clouds today





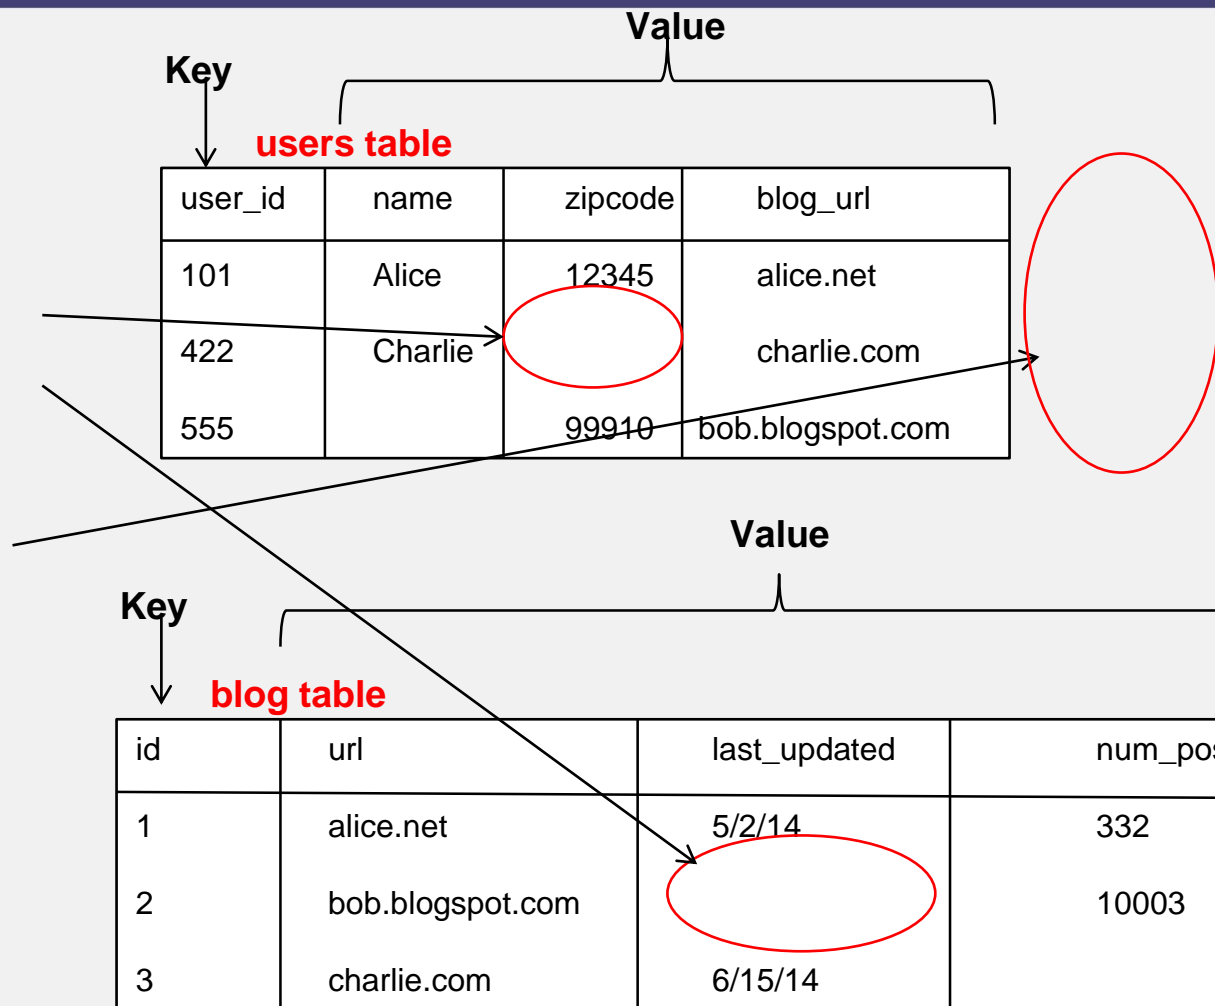
# KEY-VALUE/NoSQL DATA MODEL

- NoSQL = “Not Only SQL”
- Necessary API operations: `get(key)` and `put(key, value)`
  - And some extended operations, e.g., “CQL” in Cassandra key-value store
- Tables
  - “Column families” in Cassandra, “Table” in HBase, “Collection” in MongoDB
  - Like RDBMS tables, but ...
  - May be unstructured: May not have schemas
    - Some columns may be missing from some rows
  - Don’t always support joins or have foreign keys
  - Can have index tables, just like RDBMSs



# KEY-VALUE/NoSQL DATA MODEL

- Unstructured
- No schema imposed
- Columns missing from some rows
- No foreign keys, joins may not be supported



# COLUMN-ORIENTED STORAGE

NoSQL systems often use column-oriented storage

- RDBMSs store an entire row together (on disk or at a server)
- NoSQL systems typically store a column together (or a group of columns).
  - Entries within a column are indexed and easy to locate, given a key (and vice-versa)
- Why useful?
  - Range searches within a column are fast since you don't need to fetch the entire database
  - E.g., get me all the blog\_ids from the blog table that were updated within the past month
    - Search in the the last\_updated column, fetch corresponding blog\_id column
    - Don't need to fetch the other columns



# NEXT

Design of a real key-value store, Cassandra.

