

CLOUD COMPUTING CONCEPTS

with Indranil Gupta (Indy)

TIME AND ORDERING

Lecture A

INTRODUCTION AND BASICS

WHY SYNCHRONIZATION?

- **You want to catch a bus at 6:05 pm, but your watch is off by 15 minutes.**
 - What if your watch is late by 15 minutes?
 - You'll miss the bus!
 - What if your watch is fast by 15 minutes?
 - You'll end up unfairly waiting for a longer time than you intended.
- **Time synchronization is required for both**
 - Correctness
 - Fairness

SYNCHRONIZATION IN THE CLOUD

- Cloud airline reservation system
- Server A receives a client request to purchase last ticket on flight ABC 123.
- Server A timestamps purchase using local clock 9h:15m:32.45s, and logs it. Replies ok to client.
- That was the last seat. Server A sends message to Server B saying “flight full.”
- B enters “Flight ABC 123 full” + its own local clock value (which reads 9h:10m:10.11s) into its log.
- Server C queries A’s and B’s logs. Is confused that a client purchased a ticket at A after the flight became full at B.
- This may lead to further incorrect actions by C

WHY IS IT CHALLENGING?

- **End hosts in Internet-based systems (like clouds)**
 - Each have their own clocks
 - Unlike processors (CPUs) within one server or workstation which share a system clock
- **Processes in Internet-based systems follow an *asynchronous* system model**
 - No bounds on
 - Message delays
 - Processing delays
 - Unlike multi-processor (or parallel) systems which follow a *synchronous* system model

SOME DEFINITIONS

- An asynchronous distributed system consists of a number of [processes](#).
- Each process has a [state](#) (values of variables).
- Each process takes [actions](#) to change its state, which may be an [instruction](#) or a communication action ([send](#), [receive](#)).
- An [event](#) is the occurrence of an action.
- Each process has a local clock – events *within* a process can be assigned [timestamps](#), and thus ordered linearly.
- But – in a distributed system, we also need to know the time order of events [across](#) different processes.

CLOCK SKEW VS. CLOCK DRIFT

- Each process (running at some end host) has its own clock.
- When comparing two clocks at two processes:
 - Clock **Skew** = Relative Difference in clock *values* of two processes
 - Like distance between two vehicles on a road
 - Clock **Drift** = Relative Difference in clock *frequencies (rates)* of two processes
 - Like difference in speeds of two vehicles on the road
- A non-zero clock skew implies clocks are not synchronized.
- A non-zero clock drift causes skew to increase (eventually).
 - If faster vehicle is ahead, it will drift away
 - If faster vehicle is behind, it will catch up and then drift away

HOW OFTEN TO SYNCHRONIZE?

- Maximum Drift Rate (MDR) of a clock
- Absolute MDR is defined relative to Coordinated Universal Time (UTC). UTC is the “correct” time at any point of time.
 - MDR of a process depends on the environment.
- Max drift rate between two clocks with similar MDR is $2 * \text{MDR}$
- Given a maximum acceptable skew M between any pair of clocks, need to synchronize at least once every: $M / (2 * \text{MDR})$ time units
 - Since time = distance/speed

EXTERNAL VS INTERNAL SYNCHRONIZATION

- **Consider a group of processes**
- **External Synchronization**
 - Each process $C(i)$'s clock is within a bound D of a well-known clock S external to the group
 - $|C(i) - S| < D$ at all times
 - External clock may be connected to UTC (Universal Coordinated Time) or an atomic clock
 - E.g., Cristian's algorithm, NTP
- **Internal Synchronization**
 - Every pair of processes in group have clocks within bound D
 - $|C(i) - C(j)| < D$ at all times and for all processes i, j
 - E.g., Berkeley algorithm

EXTERNAL VS INTERNAL SYNCHRONIZATION (2)

- **External Synchronization with $D \Rightarrow$ Internal Synchronization with $2 \cdot D$**
- **Internal synchronization does not imply external synchronization**
 - In fact, the entire system may drift away from the external clock S !

NEXT

- Algorithms for clock synchronization