



CLOUD COMPUTING CONCEPTS

with Indranil Gupta (Indy)

PAXOS

Lecture B


CONSENSUS IN
SYNCHRONOUS SYSTEMS

LET'S TRY TO SOLVE CONSENSUS!

- Uh, what's the **system model**?
(assumptions!)
- **Synchronous system**: bounds on
 - Message delays
 - Upper bound on clock drift rates
 - Max time for each process stepe.g., multiprocessor (common clock across processors)
- **Processes can fail by stopping (crash-stop or crash failures)**

CONSENSUS IN SYNCHRONOUS SYSTEMS

- For a system with at most f processes crashing
 - All processes are synchronized and operate in “rounds” of time
 - the algorithm proceeds in $f+1$ rounds (with timeout), using reliable communication to all members
 - $Values_i^r$: the set of proposed values known to p_i at the beginning of round r .



CONSENSUS IN SYNCHRONOUS SYSTEM

Possible to achieve!

- For a system with at most f processes crashing
 - All processes are synchronized and operate in “rounds” of time
 - The algorithm proceeds in $f+1$ rounds (with timeout), using reliable communication to all members
 - $Values^r_i$: the set of proposed values known to p_i at the beginning of round r .
- Initially $Values^0_i = \{\}$; $Values^1_i = \{v_i\}$
 - for round = 1 to $f+1$ do
 - multicast** ($Values^r_i - Values^{r-1}_i$) // iterate through processes, send each a message
 - $Values^{r+1}_i \leftarrow Values^r_i$
 - for each V_j received
 - $Values^{r+1}_i = Values^{r+1}_i \cup V_j$
 - end
 - end
 - $d_i = \text{minimum}(Values^{f+1}_i)$

WHY DOES THE ALGORITHM WORK?

- After $f+1$ rounds, all non-faulty processes would have received the same set of values. Proof by contradiction.
- Assume that two non-faulty processes, say p_i and p_j , differ in their final set of values (i.e., after $f+1$ rounds)
- Assume that p_i possesses a value v that p_j does not possess.
 - p_i must have received v in the **very last** round
 - Else, p_i would have sent v to p_j in that last round
 - So, in the last round: a third process, p_k , must have sent v to p_i , but then crashed before sending v to p_j .
 - Similarly, a fourth process sending v in the **last-but-one round** must have crashed; otherwise, both p_k and p_j should have received v .
 - Proceeding in this way, we infer at least one (unique) crash in each of the preceding rounds.
 - This means a total of $f+1$ crashes, while we have assumed at most f crashes can occur \Rightarrow contradiction.

NEXT

- Let's be braver and solve Consensus in the Asynchronous System Model