



# CLOUD COMPUTING CONCEPTS

---

with Indranil Gupta (Indy)

## MUTUAL EXCLUSION

Lecture C

---

RICART-AGRAWALA'S ALGORITHM

# SYSTEM MODEL

- Before solving any problem, specify its System Model:
  - Each pair of processes is connected by reliable channels (such as TCP).
  - Messages are eventually delivered to recipient, and in FIFO (First In First Out) order.
  - Processes do not fail.

# RICART-AGRAWALA'S ALGORITHM

- Classical algorithm from 1981
- Invented by Glenn Ricart (NIH) and Ashok Agrawala (U. Maryland)
- No token
- Uses the notion of causality and multicast
- Has lower waiting time to enter CS than Ring-Based approach


# KEY IDEA: RICART-AGRAWALA ALGORITHM

- enter() at process  $P_i$ 
  - multicast a request to all processes
    - Request:  $\langle T, P_i \rangle$ , where  $T$  = current Lamport timestamp at  $P_i$
  - Wait until *all* other processes have responded positively to request
- Requests are granted in order of causality
- $P_i$  in request  $\langle T, P_i \rangle$  is used to break ties (since Lamport timestamps are not unique for concurrent events)


# MESSAGES IN RA ALGORITHM

- enter() at process  $P_i$ 
  - set state to Wanted
  - multicast “Request”  $\langle T_i, P_i \rangle$  to all processes, where  $T_i$  = current Lamport timestamp at  $P_i$
  - wait until all processes send back “Reply”
  - change state to Held and enter the CS
- On receipt of a Request  $\langle T_j, P_j \rangle$  at  $P_i$  ( $i \neq j$ ):
  - if (state = Held) or (state = Wanted &  $(T_i, i) < (T_j, j)$ )  
// lexicographic ordering in  $(T_j, P_j)$   
add request to local queue (of waiting requests)  
else send “Reply” to  $P_j$
- exit() at process  $P_i$ 
  - change state to Released and “Reply” to all queued requests.

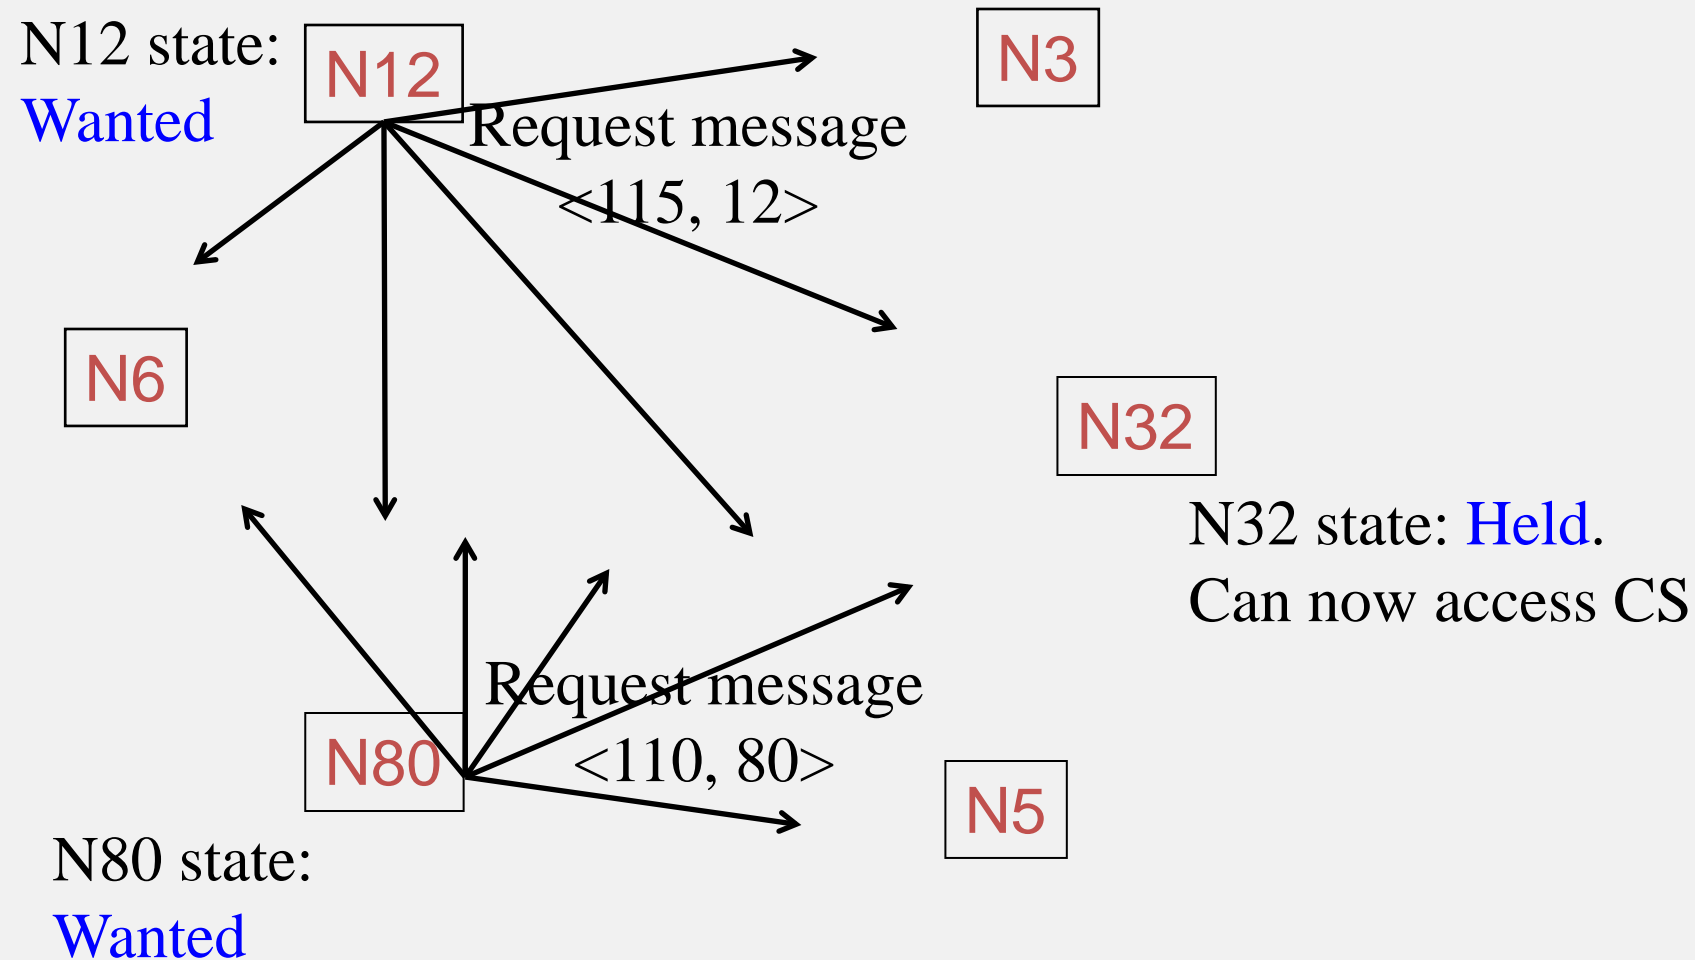
# EXAMPLE: RICART-AGRAWALA ALGORITHM



# EXAMPLE: RICART-AGRAWALA ALGORITHM

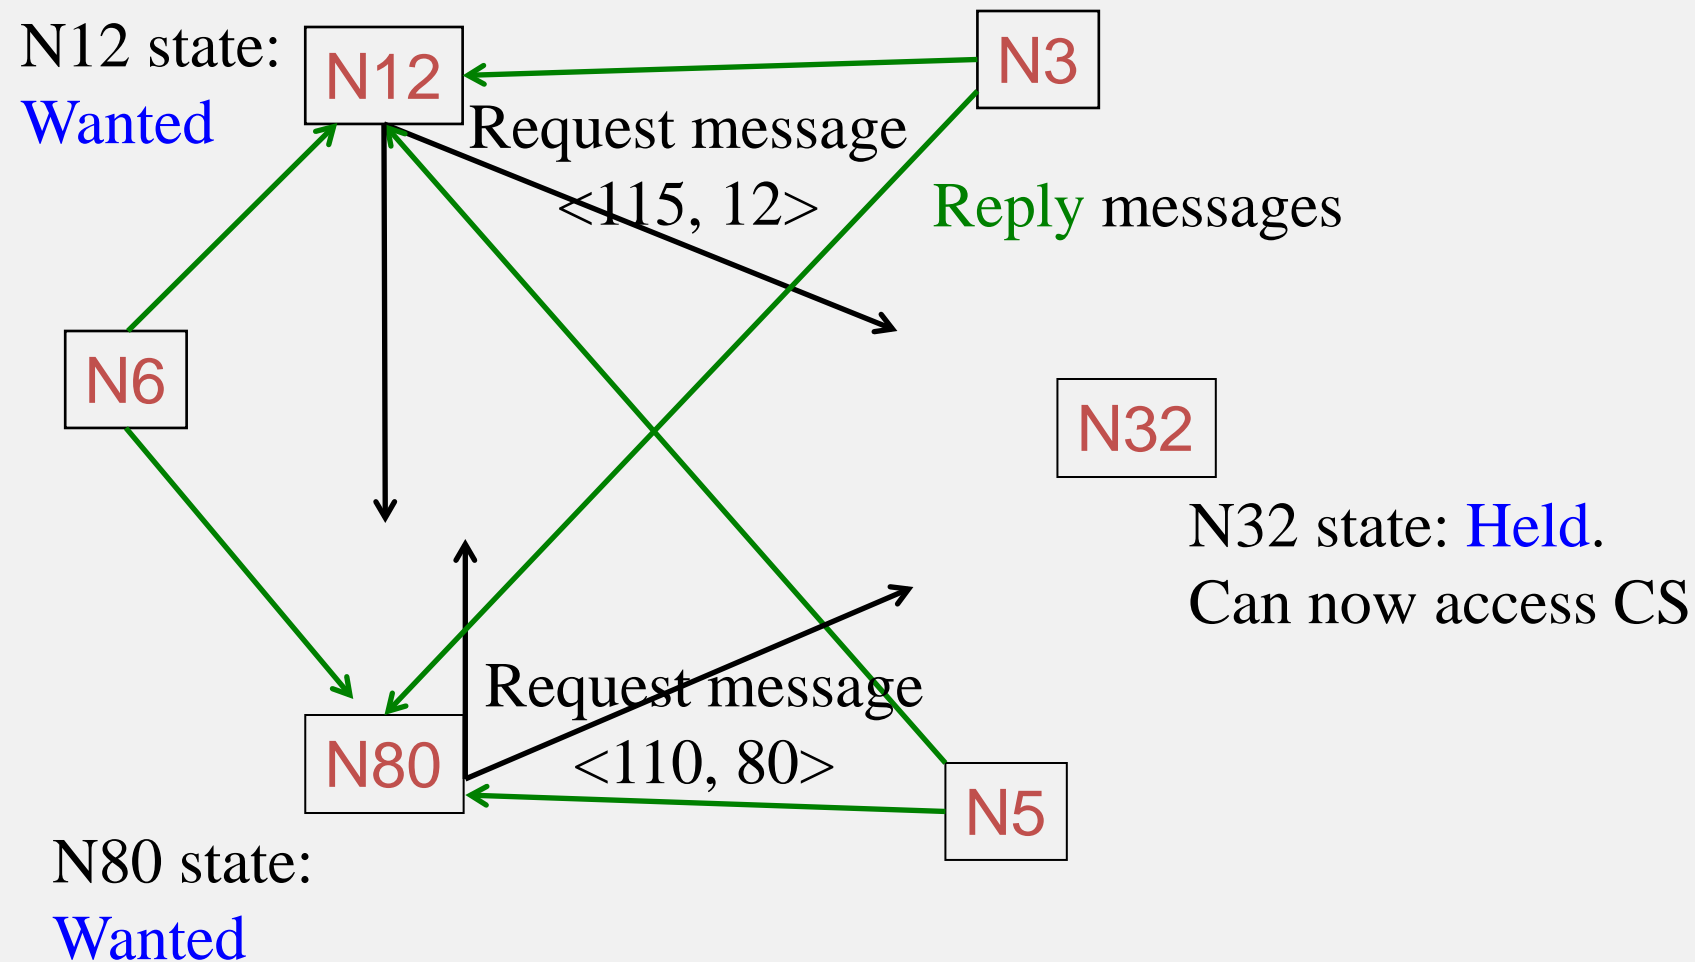


# EXAMPLE: RICART-AGRAWALA ALGORITHM

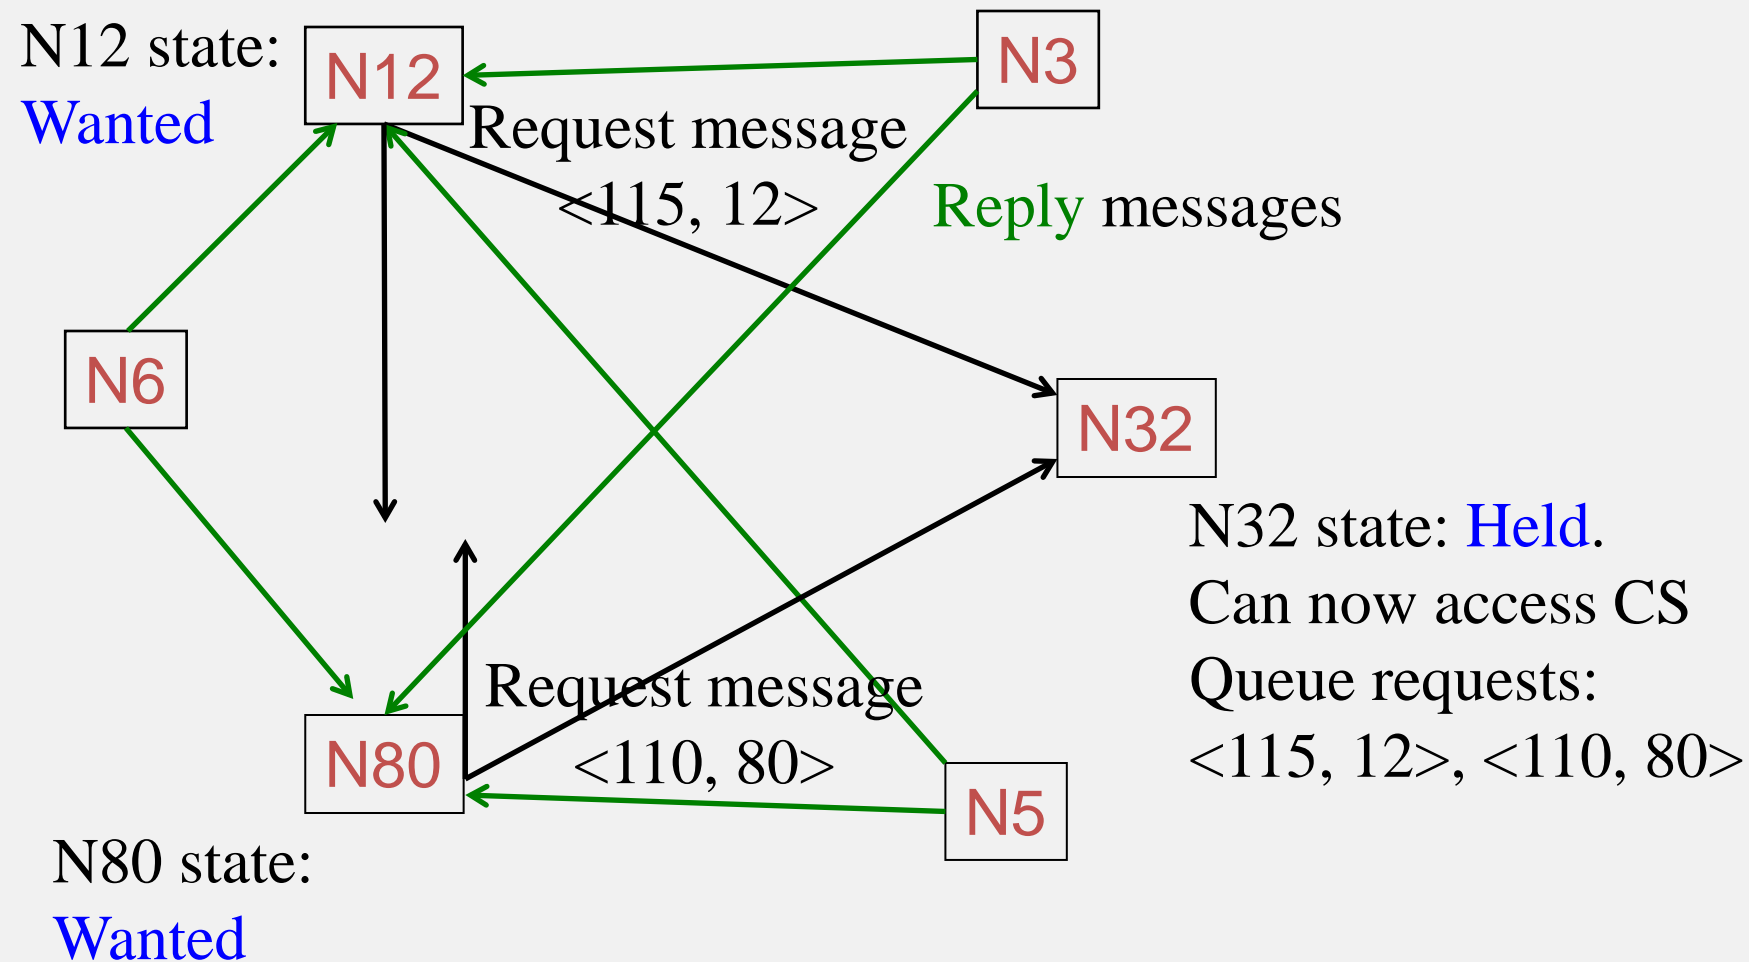




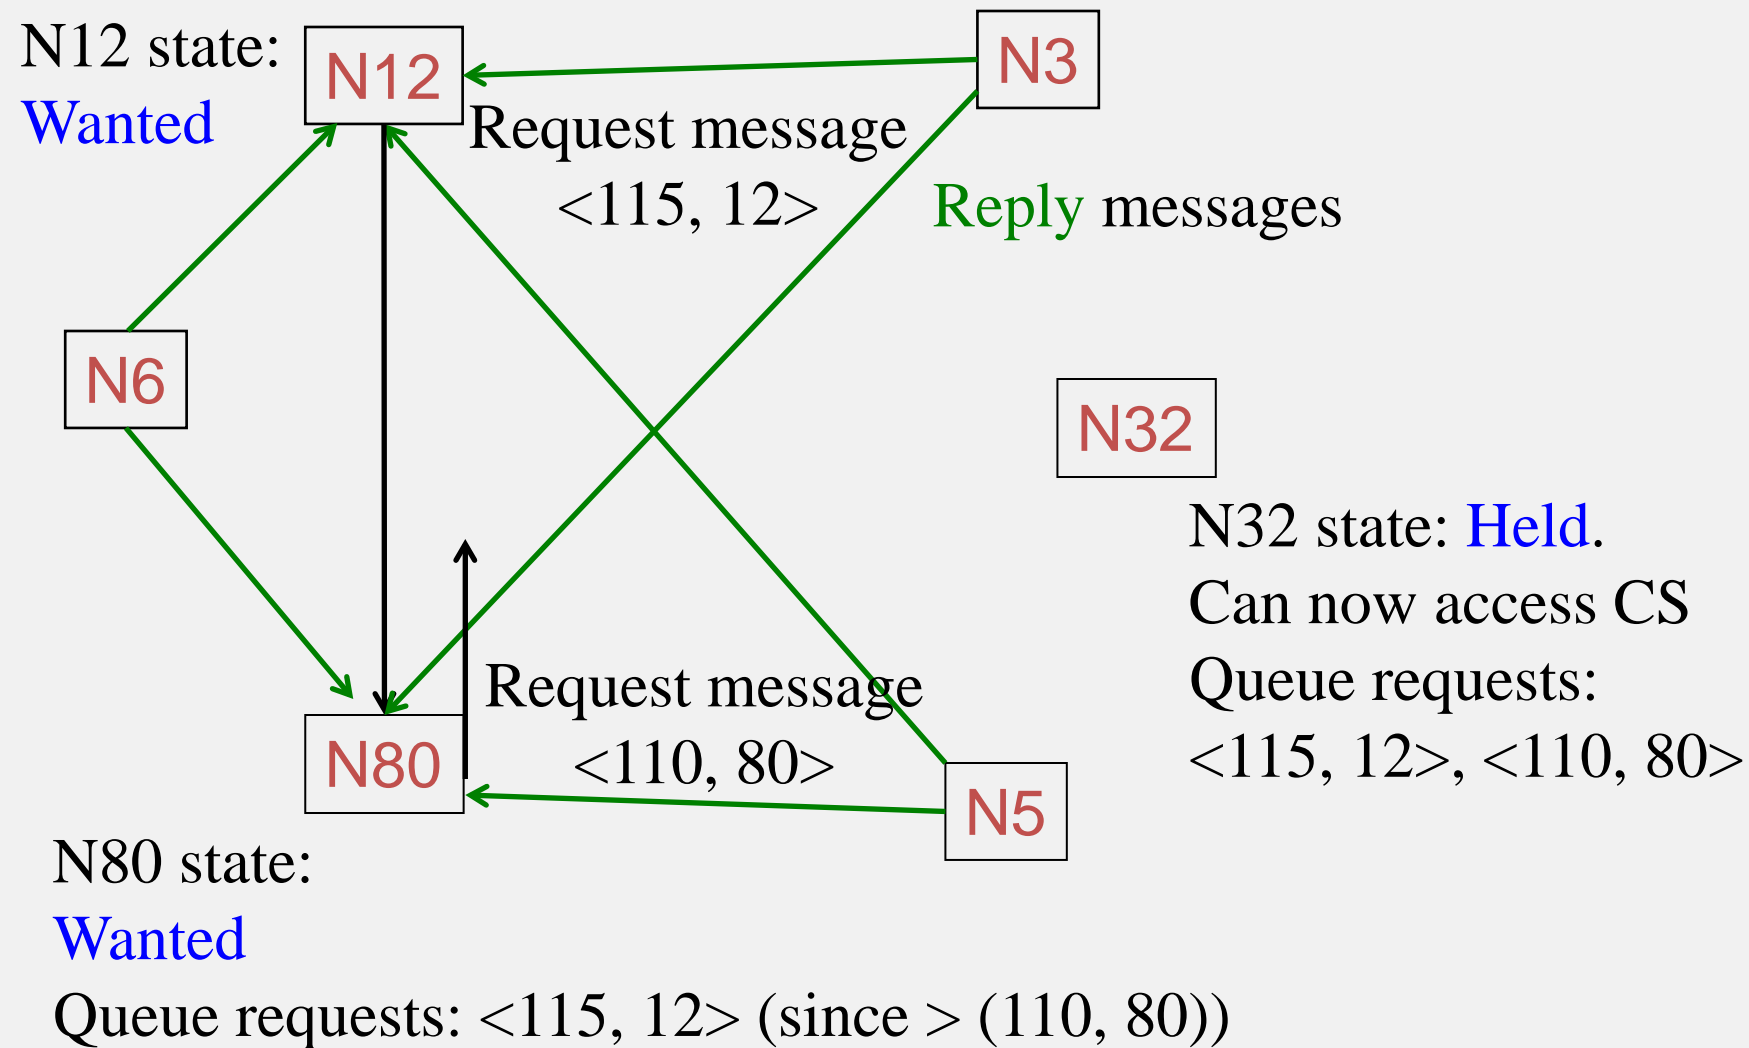
# EXAMPLE: RICART-AGRAWALA ALGORITHM



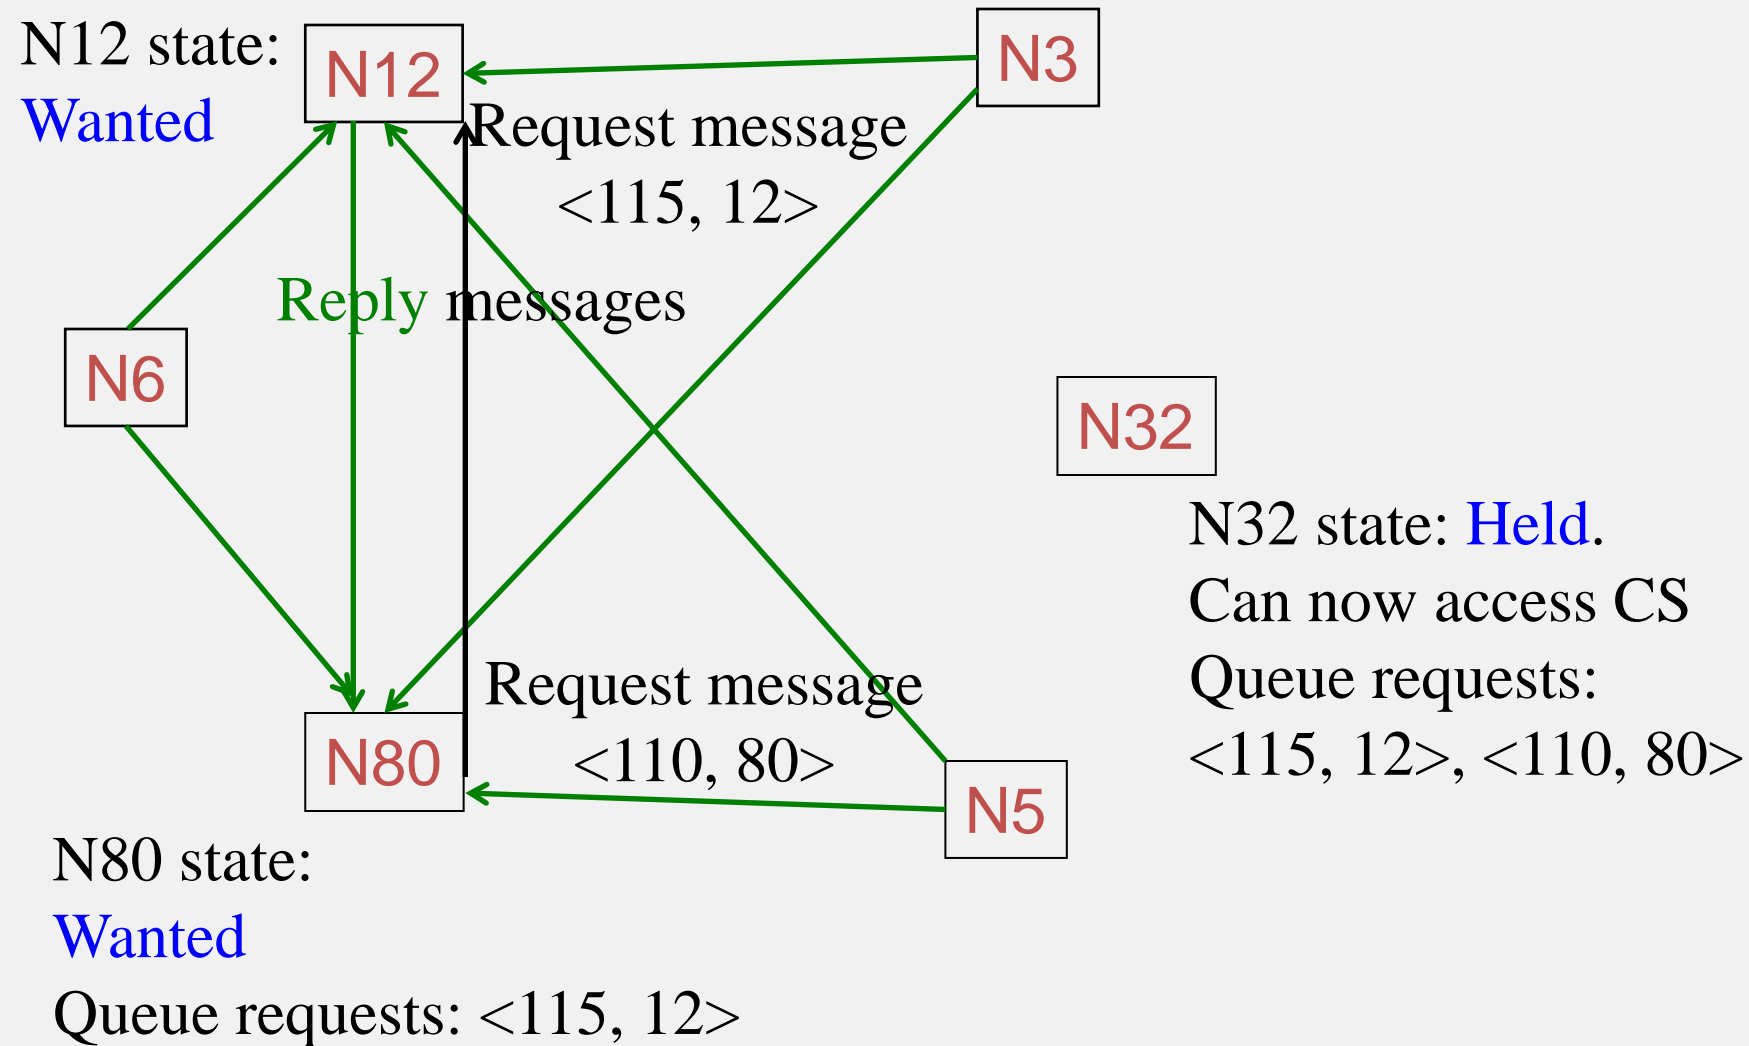
# EXAMPLE: RICART-AGRAWALA ALGORITHM



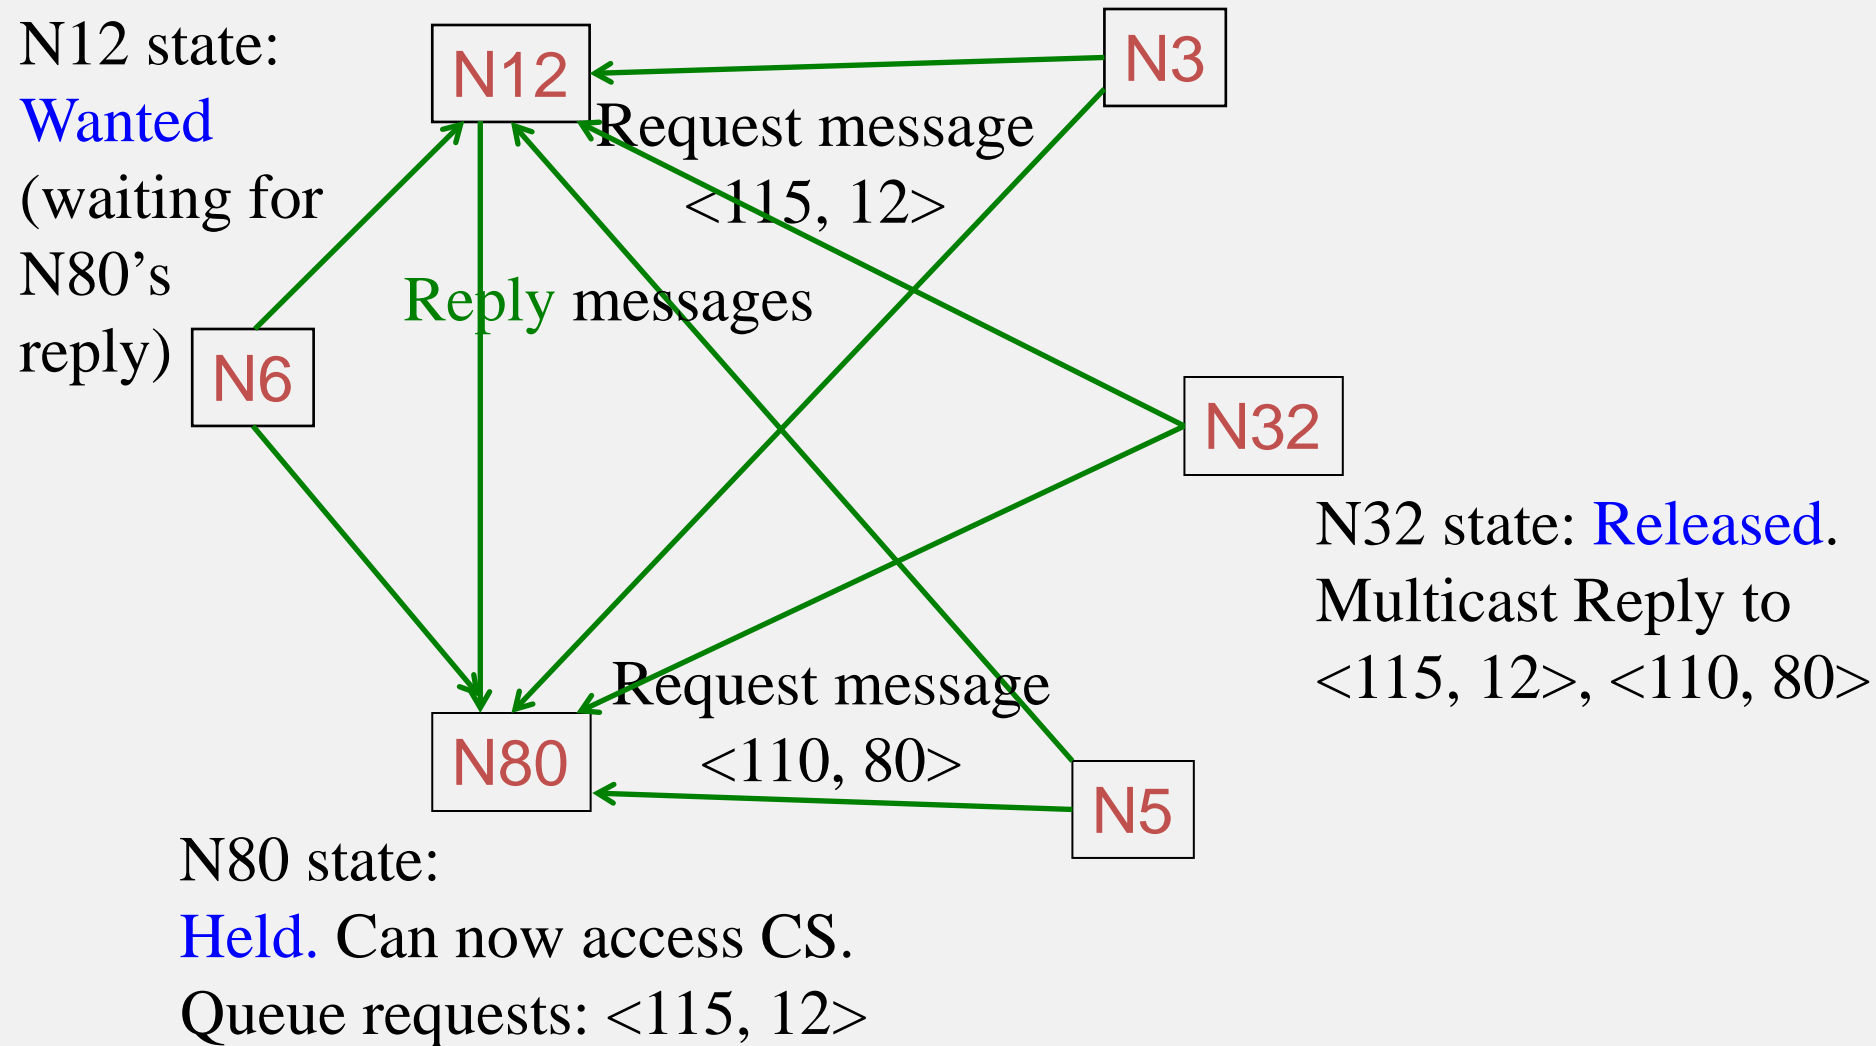
# EXAMPLE: RICART-AGRAWALA ALGORITHM



# EXAMPLE: RICART-AGRAWALA ALGORITHM



# EXAMPLE: RICART-AGRAWALA ALGORITHM



# ANALYSIS: RICART-AGRAWALA'S ALGORITHM

- Safety
  - Two processes  $P_i$  and  $P_j$  cannot both have access to CS
    - If they did, then both would have sent Reply to each other
    - Thus,  $(T_i, i) < (T_j, j)$  and  $(T_j, j) < (T_i, i)$ , which are together not possible
    - What if  $(T_i, i) < (T_j, j)$  and  $P_i$  replied to  $P_j$ 's request before it created its own request?
      - Then it seems like both  $P_i$  and  $P_j$  would approve each others' requests
      - But then, causality and Lamport timestamps at  $P_i$  implies that  $T_i > T_j$ , which is a contradiction
      - So this situation cannot arise

# ANALYSIS: RICART-AGRAWALA'S ALGORITHM (2)

- Liveness
  - Worst-case: wait for all other  $(N-1)$  processes to send Reply
- Ordering
  - Requests with lower Lamport timestamps are granted earlier

# PERFORMANCE: RICART-AGRAWALA'S ALGORITHM

- Bandwidth:  $2 \cdot (N-1)$  messages per enter() operation
  - $N-1$  unicasts for the multicast request +  $N-1$  replies
  - $N$  messages if the underlying network supports multicast
  - $N-1$  unicast messages per exit operation
    - 1 multicast if the underlying network supports multicast
- Client delay: one round-trip time
- Synchronization delay: one message transmission time



# OK, BUT ...

- Compared to Ring-Based approach, in Ricart-Agrawala approach
  - Client/synchronization delay has now gone down to  $O(1)$
  - But bandwidth has gone up to  $O(N)$
- Can we get *both* down?