

# CLOUD COMPUTING CONCEPTS with Indranil Gupta (Indy)

# MUTUAL EXCLUSION

Lecture D

MAEKAWA'S ALGORITHM AND WRAP-UP

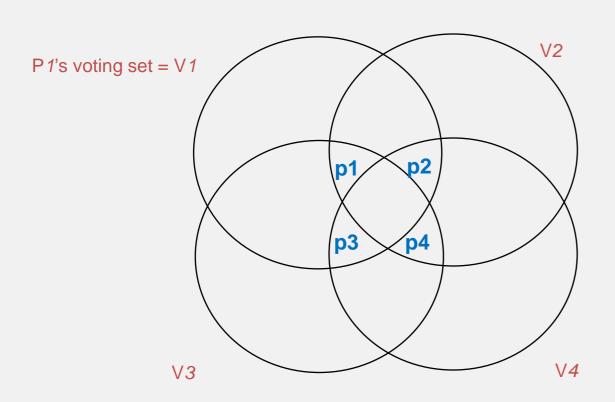
## **KEY IDEA**

- Ricart-Agrawala requires replies from *all* processes in group
- Instead, get replies from only *some* processes in group
- But ensure that only one process is given access to CS (Critical Section) at a time

#### MAEKAWA'S VOTING SETS

- Each process Pi is associated with a <u>voting set</u> Vi (of processes)
- Each process belongs to its own voting set
- The intersection of any two voting sets must be non-empty
  - Same concept as Quorums!
- Each voting set is of size *K*
- Each process belongs to *M* other voting sets
- Maekawa showed that  $K=M=\sqrt{N}$  works best
- One way of doing this is to put N processes in a  $\sqrt{N}$  by  $\sqrt{N}$  matrix and for each P*i*, its voting set V*i* = row containing P*i* + column containing P*i*. Size of voting set =  $2*\sqrt{N}-1$

# Example: Voting Sets with N=4



<b>p1</b>	<b>p2</b>
р3	p4

## Maekawa: Key Differences From Ricart-Agrawala

- Each process requests permission from only its voting set members
  - Not from all
- Each process (in a voting set) gives permission to at most one process at a time
  - Not to all

#### **ACTIONS**

- state = <u>Released</u>, voted = false
- enter() at process Pi:
  - state =  $\underline{\text{Wanted}}$
  - Multicast Request message to all processes in Vi
  - Wait for Reply (vote) messages from all processes in Vi (including vote from self)
  - state =  $\frac{\text{Held}}{}$
- exit() at process Pi:
  - state =  $\frac{\text{Released}}{\text{Released}}$
  - Multicast Reply to all processes in Vi

# Actions (2)

```
    When Pi receives a request from Pj:
    if (state == Held OR voted = true)
        queue request
    else
    send Reply to Pj and set voted = true
```

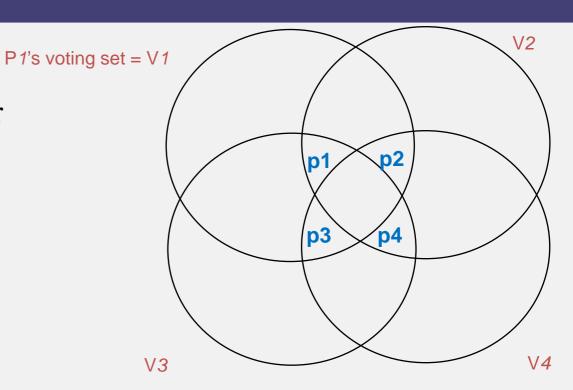
When Pi receives a Reply from Pj:
if (queue empty)
voted = false
else
dequeue head of queue, say Pk
Send Reply only to Pk
voted = true

#### SAFETY

- When a process Pi receives replies from all its voting set Vi members, no other process Pj could have received replies from all its voting set members Vj
  - Vi and Vj intersect in at least one process say Pk
  - But Pk sends only one Reply (vote) at a time, so it could not have voted for both Pi and Pj

#### LIVENESS

- A process needs to wait for at most (*N-1*) other processes to finish CS
- But does not guarantee liveness
- Since can have a *deadlock*
- Example: all 4 processes need access
  - P1 is waiting for P3
  - P3 is waiting for P4
  - P4 is waiting for P2
  - P2 is waiting for P1
  - No progress in the system!
- There are deadlock-free versions



#### **PERFORMANCE**

- Bandwidth
  - $2\sqrt{N}$  messages per enter()
  - $\sqrt{N}$  messages per exit()
  - Better than Ricart and Agrawala's (2\*(*N*-1) and *N*-1 messages)
  - $\sqrt{N}$  quite small.  $N \sim 1$  million =>  $\sqrt{N} = 1$ K
- Client delay: One round trip time
- Synchronization delay: 2 message transmission times

# WHY √N?

- Each voting set is of size *K*
- Each process belongs to *M* other voting sets
- Total number of voting set members (processes may be repeated) = K\*N
- But since each process is in *M* voting sets
  - K\*N/M = N => K = M (1)
- Consider a process Pi
  - Total number of voting sets = members present in Pi's voting set and all their voting sets = (M-1)\*K + 1
  - This must equal the number of processes
  - To minimize the overhead at each process (*K*), need each of the above members to be unique, i.e.,
    - N = (M-1)\*K + 1
    - N = (K-1)\*K + 1 (due to (1))
    - $K \sim \sqrt{N}$

### FAILURES?

- There are fault-tolerant versions of the algorithms we've discussed
  - E.g., Maekawa
- One other way to handle failures: Use Paxos-like approaches!

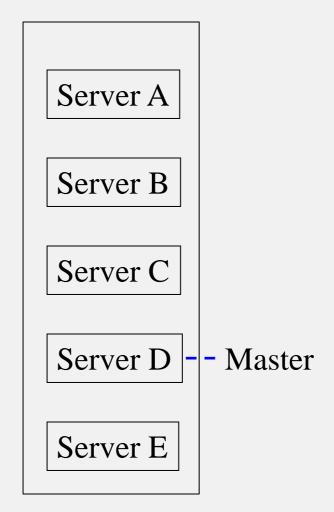
#### **CHUBBY**

- Google's system for locking
- Used underneath Google's systems like BigTable, Megastore, etc.
- Not open-sourced but published
- Chubby provides *Advisory* locks only
  - Doesn't guarantee mutual exclusion unless every client checks lock before accessing resource

Reference: http://research.google.com/archive/chubby.html

## CHUBBY (2)

- Can use not only for locking but also writing small configuration files
- Relies on Paxos
- Group of servers with one elected as Master
  - All servers replicate same information
- Clients send read requests to Master, which serves it locally
- Clients send write requests to Master, which sends it to all servers, gets majority (quorum) among servers, and then responds to client
- On master failure, run election protocol
- On replica failure, just replace it and have it catch up



#### **SUMMARY**

- Mutual exclusion important problem in cloud computing systems
- Classical algorithms
  - Central
  - Ring-based
  - Ricart-Agrawala
  - Maekawa
- Industry systems
  - Chubby: a coordination service
  - Similarly, Apache Zookeeper for coordination