



CLOUD COMPUTING CONCEPTS

with Indranil Gupta (Indy)

CONCURRENCY CONTROL

Lecture A

RPCs

WHY RPCs

- **RPC** = Remote Procedure Call
- Proposed by Birrell and Nelson in 1984
- Important abstraction for processes to call functions in other processes
- Allows code reuse
- Implemented and used in most distributed systems, including cloud computing systems
- Counterpart in Object-based settings is called RMI (Remote Method Invocation)


LOCAL PROCEDURE CALL (LPC)

- Call from one function to another function within the same process
 - Uses stack to pass arguments and return values
 - Accesses objects via pointers (e.g., C) or by reference (e.g., Java)
- LPC has *exactly-once* semantics
 - If process is alive, called function executed exactly once


REMOTE PROCEDURE CALL

- Call from one function to another function, where caller and callee function reside in different processes
 - Function call crosses a process boundary
 - Accesses objects via global references
 - Can't use pointers across processes since a reference address in process P1 may point to a different object in another process P2
 - E.g., Object address = IP + port + object number
- Similarly, RMI (Remote Method Invocation) in Object-based settings


LPCs



RPCs

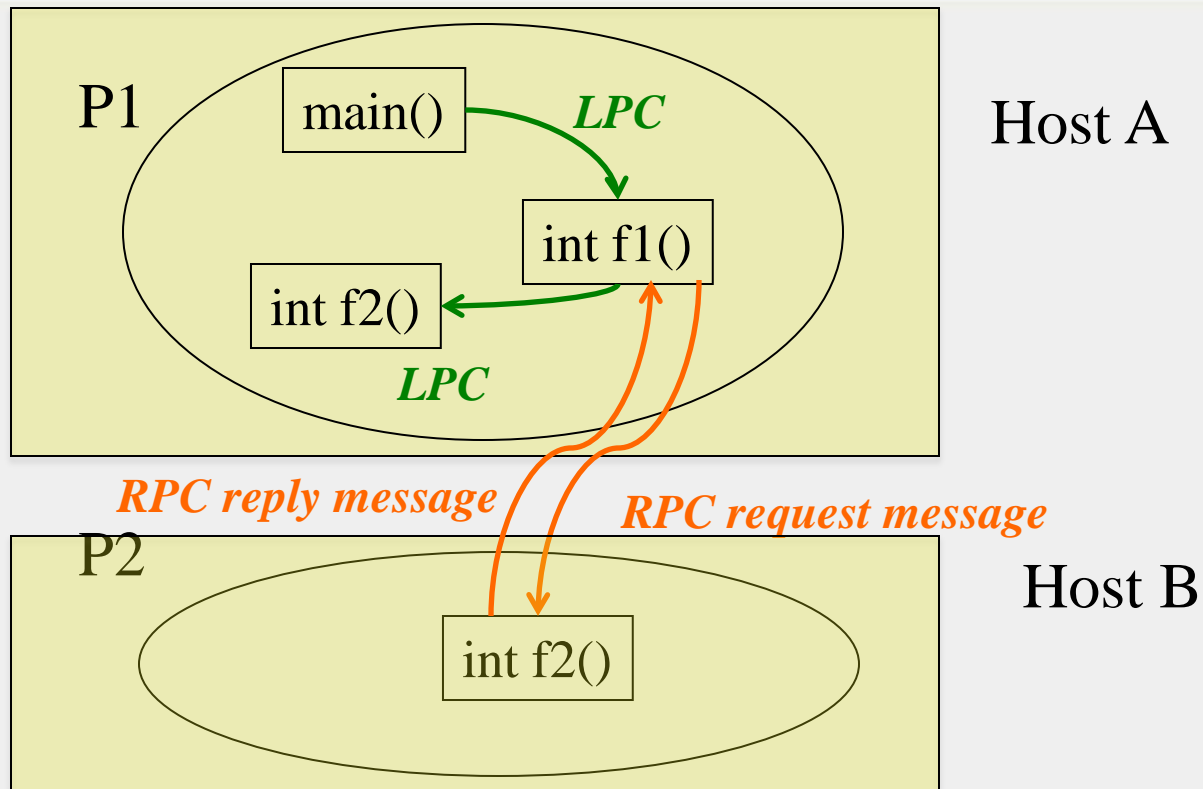


RPCs



Host A

RPCs



RPC CALL SEMANTICS

- Under failures, hard to guarantee exactly-once semantics
- Function may not be executed if
 - Request (call) message is dropped
 - Reply (return) message is dropped
 - Called process fails before executing called function
 - Called process fails after executing called function
 - Hard for caller to distinguish these cases
- Function may be executed multiple times if
 - Request (call) message is duplicated

IMPLEMENTING RPC CALL SEMANTICS

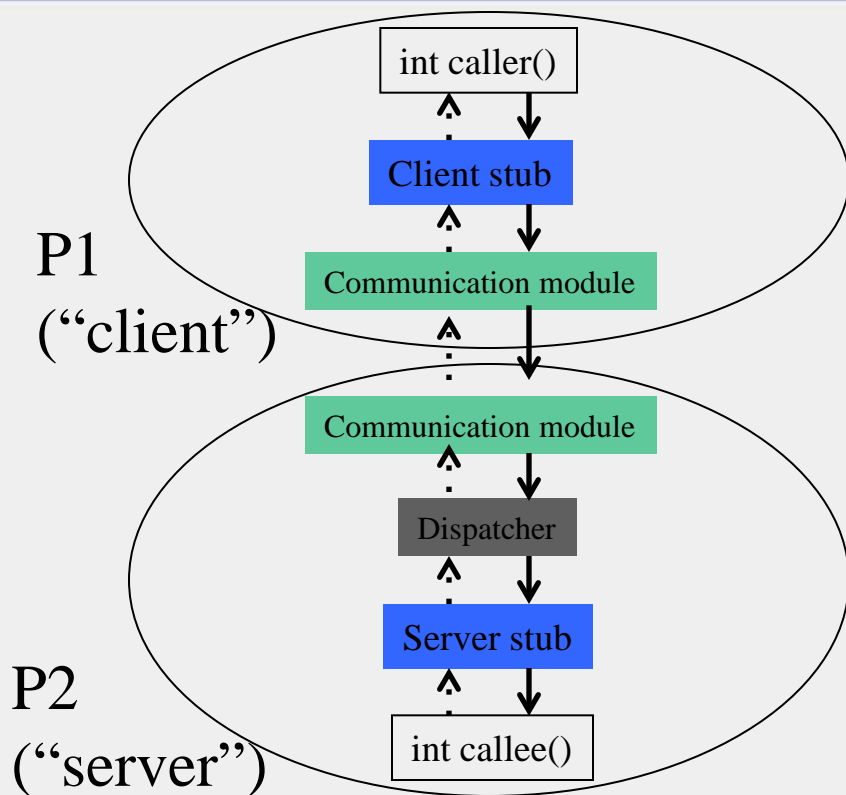
- Possible semantics
 - **At most once** semantics (e.g., Java RMI)
 - **At least once** semantics (e.g., Sun RPC)
 - Maybe, i.e., best-effort (e.g., CORBA)

Retransmit request	Filter duplicate requests	Re-execute function or retransmit reply	RPC Semantics
Yes	No	Re-execute	At least once
Yes	Yes	Retransmit	At most once
No	NA	NA	Maybe

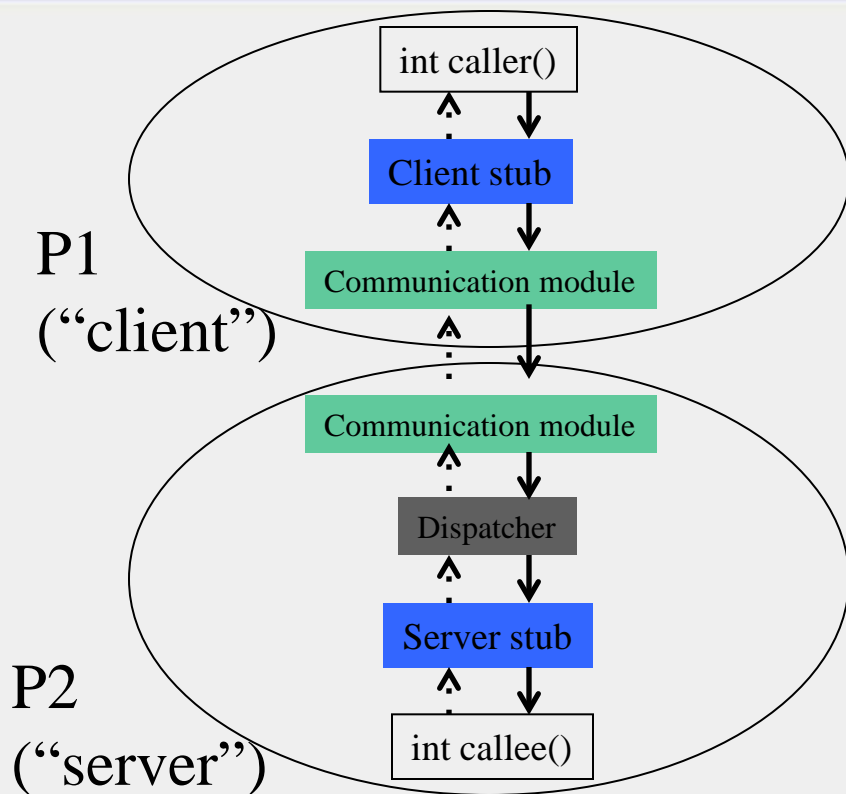
IDEMPOTENT OPERATIONS

- Idempotent operations are those that can be repeated multiple times, without any side effects
- Examples (x is server-side variable)
 - $x=1$;
 - $x=(\text{argument}) y$;
- Non-examples
 - $x=x+1$;
 - $x=x*2$;
- Idempotent operations can be used with at-least-once semantics

IMPLEMENTING RPCs



RPC COMPONENTS



Client

- **Client stub:** has same function signature as callee()
 - Allows same caller() code to be used for LPC and RPC
- **Communication Module:** Forwards requests and replies to appropriate hosts

Server

- **Dispatcher:** Selects which server stub to forward request to
- **Server stub:** calls callee(), allows it to return a value

GENERATING CODE

- Programmer only writes code for caller function and callee function
- Code for remaining components all **generated automatically** from function signatures (or object interfaces in Object-based languages)
 - E.g., Sun RPC system: Sun XDR interface representation fed into rpcgen compiler
- These components together part of a Middleware system
 - E.g., CORBA (Common Object Request Brokerage Architecture)
 - E.g., Sun RPC
 - E.g., Java RMI

MARSHALLING

- Different architectures use different ways of representing data
 - **Big endian**: Hex 12-AC-33 stored with 12 in lowest address, then AC in next higher address, then 33 in highest address
 - IBM z, System 360
 - **Little endian**: Hex 12-AC-33 stored with 33 in lowest address, then AC in next higher address, then 12
 - Intel
- Caller (and callee) process uses its own *platform-dependent* way of storing data
- Middleware has a common data representation (CDR)
 - *Platform-independent*

MARSHALLING (2)

- Middleware has a common data representation (CDR)
 - Platform-independent
- Caller process converts arguments into CDR format
 - Called “Marshalling”
- Callee process extracts arguments from message into its own platform-dependent format
 - Called “Unmarshalling”
- Return values are marshalled on callee process and unmarshalled at caller process

NEXT

- Now that we know RPCs, we can use them as a building block to understand transactions