



CLOUD COMPUTING CONCEPTS

with Indranil Gupta (Indy)

CONCURRENCY CONTROL

Lecture D

PESSIMISTIC CONCURRENCY

TWO APPROACHES

- Preventing isolation from being violated can be done in two ways
 1. **Pessimistic** concurrency control
 2. **Optimistic** concurrency control

PESSIMISTIC VS. OPTIMISTIC

- **Pessimistic**: assume the worst, prevent transactions from accessing the same object
 - E.g., Locking
- **Optimistic**: assume the best, allow transactions to write, but check later
 - E.g., Check at commit time, multi-version approaches

PESSIMISTIC: EXCLUSIVE LOCKING

- Each object has a lock
- At most one transaction can be inside lock
- Before reading or writing object O, transaction T must call **lock(O)**
 - Blocks if another transaction already inside lock
- After entering lock T can read and write O multiple times
- When done (or at commit point), T calls **unlock(O)**
 - If other transactions waiting at lock(O), allows one of them in
- Sound familiar? (This is mutual exclusion!)

CAN WE IMPROVE CONCURRENCY?

- More concurrency \Rightarrow more transactions per second \Rightarrow more revenue (\$\$\$)
- Real-life workloads have a lot of read-only or read-mostly transactions
 - Exclusive locking reduces concurrency
 - Hint: Ok to allow two transactions to concurrently read an object, since read-read is not a conflicting pair

ANOTHER APPROACH: READ-WRITE LOCKS

- Each object has a lock that can be held in one of two modes
 - **Read mode**: multiple transactions allowed in
 - **Write mode**: exclusive lock
- Before first reading O, transaction T calls `read_lock(O)`
 - T allowed in only if *all* transactions inside lock for O all entered via read mode
 - Not allowed if *any* transaction inside lock for O entered via write mode

READ-WRITE LOCKS (2)

- Before first writing O, call `write_lock(O)`
 - Allowed in only if no other transaction inside lock
- If T already holds `read_lock(O)`, and wants to write, call `write_lock(O)` to *promote* lock from read to write mode
 - Succeeds only if no other transactions in write mode or read mode
 - Otherwise, T blocks
- `Unlock(O)` called by transaction T releases any lock on O by T

GUARANTEERING SERIAL EQUIVALENCE WITH LOCKS

- Two-phase locking

- A transaction cannot acquire (or promote) any locks after it has started releasing locks
- Transaction has two phases
 1. Growing phase: only acquires or promotes locks
 2. Shrinking phase: only releases locks
 - Strict two phase locking: releases locks only at commit point

WHY TWO-PHASE LOCKING \Rightarrow SERIAL EQUIVALENCE?



- Proof by contradiction
- Assume two-phase locking system where serial equivalence is violated for some two transactions T1, T2
- Two facts must then be true:
 - (A) For some object O1, there were conflicting operations in T1 and T2 such that the time ordering pair is (T1, T2)
 - (B) For some object O2, the conflicting operation pair is (T2, T1)
- (A) \Rightarrow T1 released O1's lock and T2 acquired it after that
 \Rightarrow T1's shrinking phase is before or overlaps with T2's growing phase
- Similarly, (B) \Rightarrow T2's shrinking phase is before or overlaps with T1's growing phase
- But both these cannot be true!

DOWNSIDE OF LOCKING

- Deadlocks!

DOWNSIDE OF LOCKING – DEADLOCKS!

Transaction T1

Lock(ABC123);

x = write(10, ABC123);

Lock(ABC789);

// Blocks waiting for T2

...


Transaction T2

Lock(ABC789);

y = write(15, ABC789);

Lock(ABC123);

... // Blocks waiting for T1



WHEN DO DEADLOCKS OCCUR?

- 3 necessary conditions for a deadlock to occur
 1. Some objects are accessed in exclusive lock modes
 2. Transactions holding locks cannot be preempted
 3. There is a circular wait (cycle) in the Wait-for graph
- “Necessary” = if there’s a deadlock, these conditions are all definitely true
- (Conditions not sufficient: if they’re present, it doesn’t imply a deadlock is present.)

COMBATING DEADLOCKS

1. Lock **timeout**: abort transaction if lock cannot be acquired within timeout
 - ☹ Expensive, wasted work
2. Deadlock **Detection**:
 - Keep track of Wait-for graph (e.g., via Global Snapshot algorithm), and
 - Find cycles in it (e.g., periodically)
 - If find cycle, there's a deadlock => Abort one or more transactions to break cycle
 - ☹ Still allows deadlocks to occur

COMBATING DEADLOCKS (2)

3. Deadlock Prevention

- Set up the system so one of the *necessary conditions* is violated
 1. *Some objects are accessed in exclusive lock modes*
 - Fix: Allow read-only access to objects
 2. *Transactions holding locks cannot be preempted*
 - Fix: Allow preemption of some transactions
 3. *There is a circular wait (cycle) in the Wait-for graph*
 - Fix: Lock all objects in the beginning; if fail any, abort transaction
=> No cycles in Wait-for graph

NEXT

- Can we allow more concurrency?
- Optimistic Concurrency Control