



CLOUD COMPUTING CONCEPTS

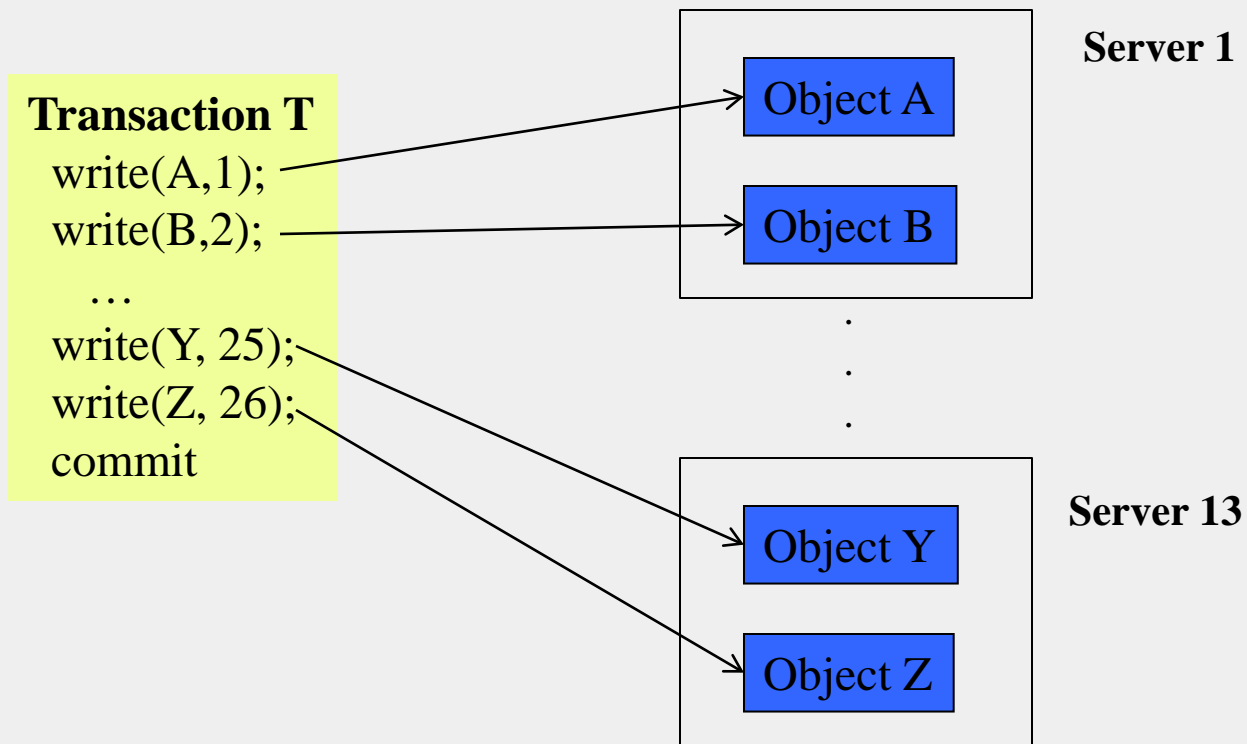
with Indranil Gupta (Indy)

REPLICATION CONTROL

Lecture B

TWO-PHASE COMMIT

TRANSACTIONS WITH DISTRIBUTED SERVERS



TRANSACTIONS WITH DISTRIBUTED SERVERS (2)

- Transaction T may touch objects that reside on different servers
- When T tries to commit
 - Need to ensure all these servers commit their updates from T \Rightarrow T will commit
 - Or none of these servers commit \Rightarrow T will abort
- What problem is this?


TRANSACTIONS WITH DISTRIBUTED SERVERS (2)

- Transaction T may touch objects that reside on different servers
- When T tries to commit
 - Need to ensure all these servers commit their updates from T \Rightarrow T will commit
 - Or none of these servers commit \Rightarrow T will abort
- What problem is this?
 - Consensus!
 - (It's called the "Atomic Commit problem")

ONE-PHASE COMMIT

Transaction T

```
write(A,1);  
write(B,2);  
...  
write(Y, 25);  
write(Z, 26);  
commit
```




- *Special server called “Coordinator” initiates atomic commit*
- *Tells other servers to either commit or abort*


ONE-PHASE COMMIT: ISSUES

- Server with object has no say in whether transaction commits or aborts
 - If object corrupted, it cannot commit (while other servers have committed)
- Server may crash before receiving commit message, with some updates still in memory


TWO-PHASE COMMIT




TWO-PHASE COMMIT



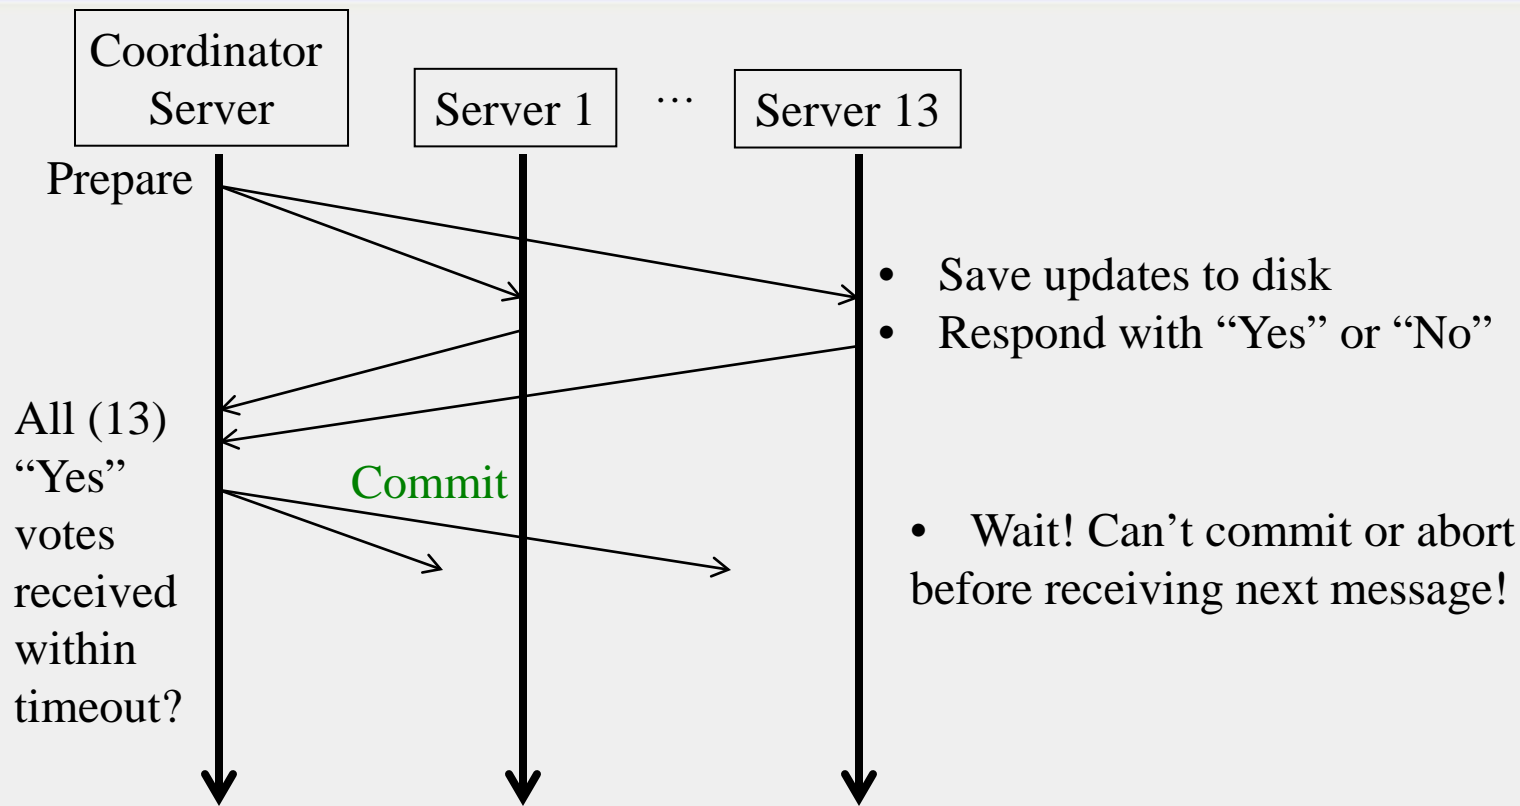
TWO-PHASE COMMIT



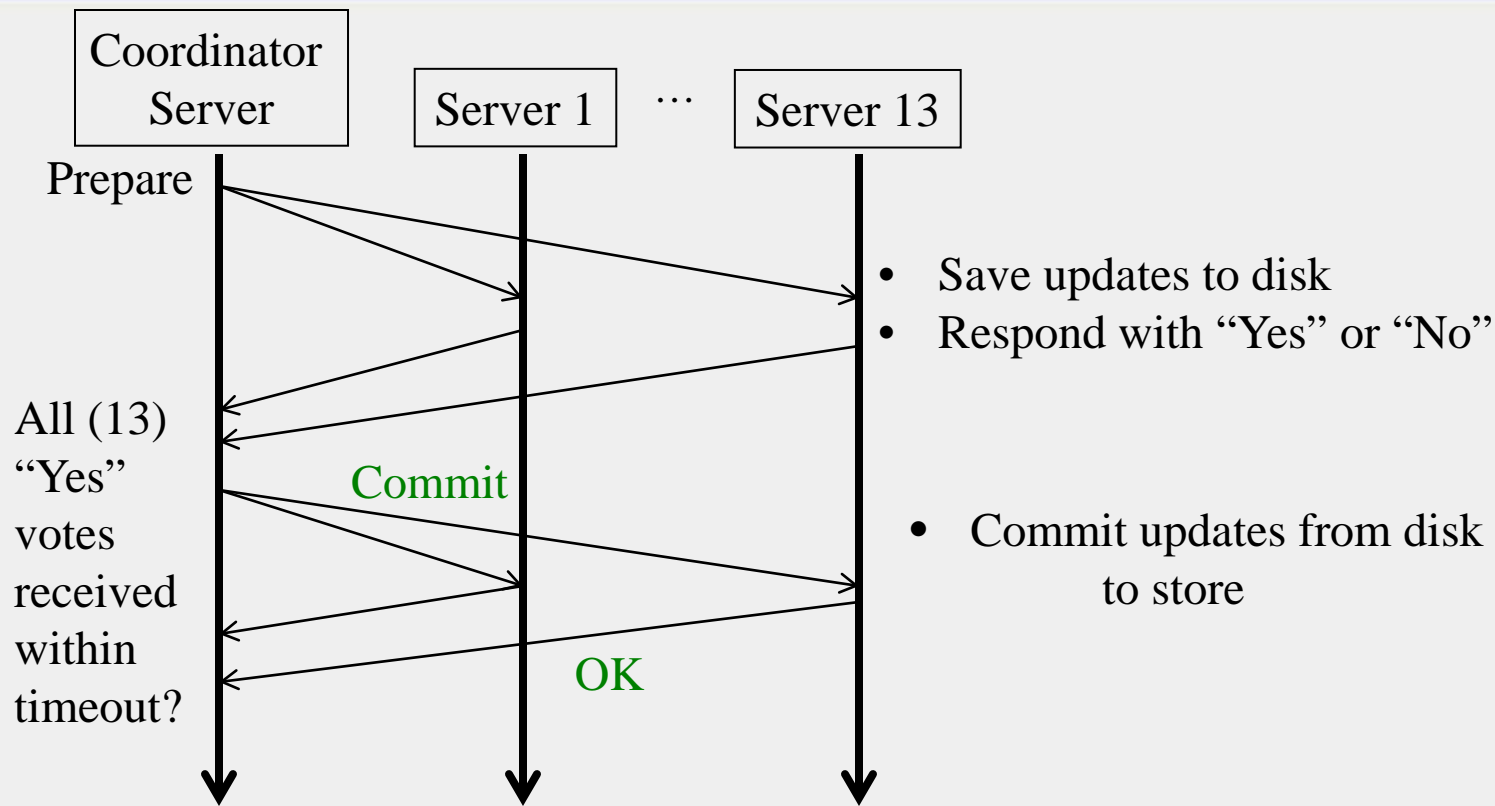
TWO-PHASE COMMIT



TWO-PHASE COMMIT



TWO-PHASE COMMIT



FAILURES IN TWO-PHASE COMMIT

- To deal with server crashes
 - Each server saves tentative updates into permanent storage, right before replying Yes/No in first phase. Retrievable after crash recovery.
 - Coordinator logs votes and decisions too
- To deal with coordinator crashes
 - Coordinator logs all decisions and received/sent messages on disk
 - After recovery or new election => new coordinator takes over
- To deal with Prepare message loss
 - The server may decide to abort unilaterally after a timeout for first phase (server will always vote No, and so coordinator will also eventually abort)

FAILURES IN TWO-PHASE COMMIT (2)

- To deal with Yes/No message loss, coordinator aborts the transaction after a timeout (pessimistic!). It must announce Abort message to all.
- To deal with Commit or Abort message loss
 - Server can poll coordinator (repeatedly)
- If server voted Yes, it cannot commit unilaterally before receiving Commit message
- If server voted No, can abort right away (why?)

USING PAXOS IN DISTRIBUTED SERVERS

Atomic Commit

- Can also use Paxos to decide whether to commit a transaction or not
- But need to ensure that if any server votes No, everyone aborts

Ordering updates

- Paxos can also be used by replica group (for an object) to order all updates
 - Server proposes message for next sequence number
 - Group reaches consensus (or not)

SUMMARY

- Multiple servers in cloud
 - Replication for fault-tolerance
 - Load balancing across objects
- Replication flavors using concepts we learnt earlier
 - Active replication
 - Passive replication
- Transactions and distributed servers
 - Two-phase commit