



CLOUD COMPUTING CONCEPTS

with Indranil Gupta (Indy)

STREAM PROCESSING

Lecture A

STREAM PROCESSING IN STORM

WHAT WE'LL COVER

- Why Stream Processing
- Storm

STREAM PROCESSING CHALLENGE

- Large amounts of data => Need for real-time views of data
 - Social network trends, e.g., Twitter real-time search
 - Website statistics, e.g., Google Analytics
 - Intrusion detection systems, e.g., in most datacenters
- Process large amounts of data
 - With latencies of few seconds
 - With high throughput

MAPREDUCE?

- Batch Processing => Need to wait for entire computation on large dataset to complete
- Not intended for long-running stream-processing

ENTER STORM

- Apache Project
- <https://storm.incubator.apache.org/>
- Highly active JVM project
- Multiple languages supported via API
 - Python, Ruby, etc.
- Used by over 30 companies including
 - Twitter: For personalization, search
 - Flipboard: For generating custom feeds
 - Weather Channel, WebMD, etc.

STORM COMPONENTS

- Tuples
- Streams
- Spouts
- Bolts
- Topologies

TUPLE

- An ordered list of elements
- E.g., < tweeter, tweet >
 - E.g., < “Miley Cyrus”, “Hey! Here’s my new song!” >
 - E.g., < “Justin Bieber”, “Hey! Here’s MY new song!” >
- E.g., < URL, clicker-IP, date, time >
 - E.g., < coursera.org, 101.201.301.401, 4/4/2014, 10:35:40 >
 - E.g., < coursera.org, 901.801.701.601, 4/4/2014, 10:35:42 >

Tuple


STREAM

- Sequence of tuples
 - Potentially unbounded in number of tuples
- Social network example:
 - $\langle \text{"Miley Cyrus"}, \text{"Hey! Here's my new song!"} \rangle$,
 - $\langle \text{"Justin Bieber"}, \text{"Hey! Here's MY new song!"} \rangle$,
 - $\langle \text{"Rolling Stones"}, \text{"Hey! Here's my old song that's still a super-hit!"} \rangle$, ...
- Website example:
 - $\langle \text{coursera.org}, 101.201.301.401, 4/4/2014, 10:35:40 \rangle$,
 - $\langle \text{coursera.org}, 901.801.701.601, 4/4/2014, 10:35:42 \rangle$, ...



SPOUT


- A Storm entity (process) that is a source of streams
- Often reads from a crawler or DB



BOLT




- A Storm entity (process) that
 - Processes input streams
 - Outputs more streams for other bolts




TOPOLOGY

- A directed graph of spouts and bolts (and output bolts)
- Corresponds to a Storm “application”



TOPOLOGY

- Can have cycles if the application requires it



BOLTS COME IN MANY FLAVORS

- Operations that can be performed
 - **Filter:** forward only tuples which satisfy a condition
 - **Joins:** When receiving two streams A and B, output all pairs (A,B) which satisfy a condition
 - **Apply/transform:** Modify each tuple according to a function
 - And many others
- But bolts need to process a lot of data
 - Need to make them fast

PARALLELIZING BOLTS

- Have multiple processes (“tasks”) constitute a bolt
- Incoming streams split among the tasks
- Typically each incoming tuple goes to one task in the bolt
 - Decided by “**Grouping strategy**”
- Three types of grouping are popular

GROUPING

- **Shuffle Grouping**
 - Streams are distributed evenly across the bolt's tasks
 - Round-robin fashion
- **Fields Grouping**
 - Group a stream by a subset of its fields
 - E.g., All tweets where twitter username starts with [A-M,a-m,0-4] goes to task 1, and all tweets starting with [N-Z,n-z,5-9] go to task 2
- **All Grouping**
 - All tasks of bolt receive all input tuples
 - Useful for joins

STORM CLUSTER

- Master node
 - Runs a daemon called *Nimbus*
 - Responsible for
 - Distributing code around cluster
 - Assigning tasks to machines
 - Monitoring for failures of machines
- Worker node
 - Runs on a machine (server)
 - Runs a daemon called *Supervisor*
 - Listens for work assigned to its machines
- Zookeeper
 - Coordinates Nimbus and Supervisors communication
 - All state of Supervisor and Nimbus is kept here

FAILURES

- A tuple is considered failed when its topology (graph) of resulting tuples fails to be fully processed within a specified timeout
- **Anchoring:** Anchor an output to one or more input tuples
 - Failure of one tuple causes one or more tuples to replayed

API FOR FAULT-TOLERANCE (OUTPUTCOLLECTOR)

- **Emit**(tuple, output)
 - Emits an output tuple, perhaps anchored on an input tuple (first argument)
- **Ack**(tuple)
 - Acknowledge that you finish processing a tuple
- **Fail**(tuple)
 - Immediately fail the spout tuple at the root of tuple topology if there is an exception from the database, etc.
- Must remember to ack/fail each tuple
 - Each tuple consumes memory. Failure to do so results in memory leaks.

SUMMARY

- Processing data in real-time a big requirement today
- Storm
 - And other sister systems, e.g., Spark Streaming
- Parallelism
- Application topologies
- Fault-tolerance